

Table of Contents

- [Table of Contents](#)
 - [Kubernetes Terminology](#)
 - [Kubernetes Architecture](#)
 - [Kubernetes Interactive Labs](#)
 - [Kubernetes Objects](#)
 - [Pods](#)
 - [ReplicaSets](#)
 - [Scaling Replicaset](#)
 - [Deployments](#)
 - [Namespaces](#)
 - [Kubernetes Namespaces Concepts](#)
 - [Kubernetes Service](#)
 - [NodePort](#)
 - [ClusterIP](#)
 - [LoadBalancer](#)
 - [Kubernetes Service vs Deployment](#)
 - [Reference](#)
 - [Kubernetes Cluster Setup](#)
-

Kubernetes Terminology

- **Nodes:** Hosts that run Kubernetes applications/containers
- **Containers:** Units of packaging application
- **Pods:** Units of deployment which is collection of containers
- **Replication Controller:** Ensures availability and scalability
- **Labels:** Key-value pairs for identification (similar to tags)
- **Services:** Collection of pods exposed as an endpoint

--

Kubernetes Architecture

- **Master** controls the cluster, and the nodes in it. It ensures the execution only happens in nodes and co-ordinates the container execution on nodes.
- **Nodes** host the containers; in-fact these Containers are grouped logically to form **Pods**.
- Each node can run multiple such **Pods**, which are a group of containers, that interact with each other, for a **deployment**.
- **Replication Controller** is Master's resource to ensure that the requested no. of pods are always running on nodes.
- **Replica Set:** replica sets are created by deployment, these deployments contains declaration of containers which you want to run in cluster like image/tag, env variable, data volumes Kubernetes has several components in its architecture.
- **Service** is an object on Master that provides load balancing across a replicated group of Pods.

- **POD**: POD is the smallest unit of deployment in K8S object. Each POD then contains the Docker containers. Each POD can host a different set of Docker containers. The proxy is then used to control the exposing of these services to the outside world. PODs are created by replicaset.

--

- **etcd** – This component is a highly available **key-value** store that is used for storing shared configuration and service discovery.
- **kube-apiserver** – This is an API which can be used to orchestrate the Docker containers.
- **kube-controller-manager** – This is used to control the Kubernetes services.
- **kube-scheduler** – This is used to schedule the containers on hosts.
- **Kubelet** – This is used to control the launching of containers via manifest files from worker host. (which talks with K8S cluster).
- **kube-proxy** – This is used to provide network proxy services to the outside world.

Kubernetes Interactive Labs

- Create an account with DockerHub and Sign-In
- Navigate to [Play with Kubernetes](#)
- Click on **Add New Instance** on the left side of the screen to bring up an instance on the right side.
- Perform the steps mentioned in the terminal.
 - Copy the commands using **Ctrl+Insert** and Paste the commands using **Shift+Insert**
 - To maximize the window use **Alt+Enter**

```
# Initializes cluster master node:
kubeadm init --apiserver-advertise-address $(hostname -i) --pod-network-cidr
10.5.0.0/16

# Initialize cluster networking:
kubectl apply -f https://raw.githubusercontent.com/cloudnativelabs/kube-
router/master/daemonset/kubeadm-kuberouter.yaml

# To get node information
kubectl get nodes

# (Optional) Create an nginx deployment:
# kubectl apply -f
https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/ap
plication/nginx-app.yaml
```

--

- To add Worker Node: Click on **Add New Instance**
- Copy the output of 1st command from 1st instance into the New Instance

```
kubeadm join 192.168.0.13:6443 --token n8lvoe.2h1ubmjyfmcxr00s \
--discovery-token-ca-cert-hash
```

```
sha256:089afdc26c7908fd5d0ec54d9e812a85ac6dc194457968961ccaa19d8f7cb7a3
```

- Execute the commands on the Control Plane Node.

```
kubectl get nodes
```

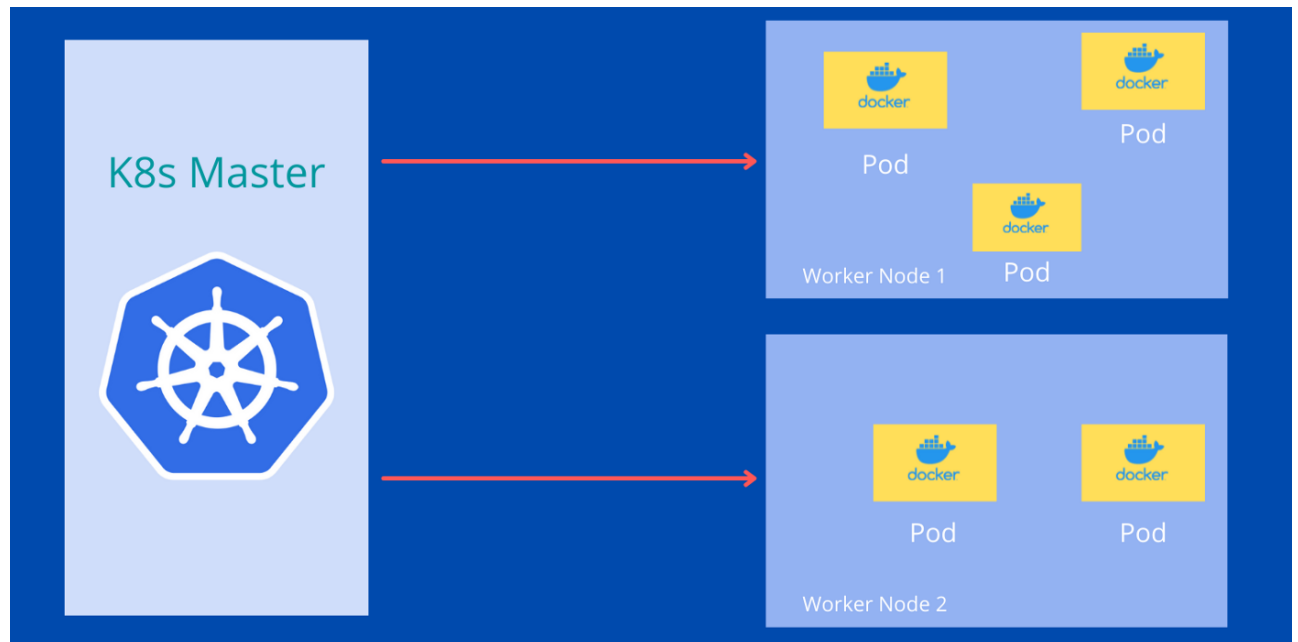
Kubernetes Objects

- Pods
- ReplicaSets
- Deployments
- Namespaces
- Service
- Configmap
- Secrets

Pods

- Kubernetes doesn't deploy containers directly on the worker node.
- Kubernetes is designed to handle single or multiple containers. Containers are encapsulated into a Kubernetes object. That Kubernetes object is called **Pod**.
 - If we have to run a Docker container in the Kubernetes cluster then we need a Pod object.
- A Pod is the smallest and simplest Kubernetes object.
- The **Pod** is one of the Kubernetes components and the smallest unit in the Kubernetes cluster.
- A group of one or more containers is called a Pod.
- Containers in a Pod are deployed together, and are started, stopped, and replicated as a group.
- Each Pod has only 1 IP, irrespective of number of containers.
- All container in a Pod shares IP, cgroups, namespaces, localhost adapter, volumes every pod can interact directly with other pod via Pod N/W (Inter-Pod communication)
- It is the unit of deployment in Kubernetes, which represents a single instance of the application.

--



•

--

- Create a Pod with **Imperative Command**.

```
# Imperative Command is simply a kubectl command line with options to create
objects.
# Syntax
kubectl run pod-name --image image_name:tag

kubectl run nginx-pod --image nginx:1.16.2
# kubectl run => Standard command to run a Pod.
# nginx-pod => Name of the Pod. You can give any name.
# --image nginx:1.16.2 => This is the Docker image that will be used to launch the
Container in the Pod.
kubectl get pods

kubectl get pods -o wide

kubectl describe pod nginx-pod

kubectl run new-nginx-pod --image nginx:1.16.1

# View Pods in the cluster
kubectl get pods

# Check for details of the pod
kubectl describe pod new-nginx-pod
```

--

- Create a Pod with **Declarative Command**, where Kubernetes objects can be created, updated, and deleted by storing multiple object configuration files in a directory and using kubectl apply to recursively create and update those objects as needed.

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: testapp-pod
  labels:
    app: testapp-pod
    type: front-end-pod
spec:
  containers:
  - name: nginx-container
    image: nginx:1.16.1
```

- Execute below commands to apply and validate the YAML manifest file.

```
kubectl get pods -o wide
```

```
kubectl create -f pod-definition.yaml
```

OR

```
kubectl apply -f pod-definition.yaml
```

Get the list of running Pods using the below command

```
kubectl get pods
```

OUTPUT

NAME	READY	STATUS	RESTARTS	AGE
testapp-pod	1/1	Running	0	10s

Check for details of the pod

```
kubectl describe pod testapp-pod
```

List pods with wide option which also provide Node information on which Pod is running.

```
kubectl get pods -o wide
```

Login inside the pod using below command

```
kubectl exec -it testapp-pod /bin/bash
```

In your shell, list the root directory:

```
root@testapp-pod:/# ls /
```

```
bin    dev                                docker-entrypoint.sh  home  lib64  mnt  proc  run   srv
tmp    var
boot   docker-entrypoint.d  etc                   lib   media  opt  root  sbin  sys
usr
```

You can run these example commands inside the container

```
ls /
```

```
apt-get update
```

```

apt-get install -y procps curl
ps aux
ps aux | grep nginx

# Modifying the root page for nginx Webserver
curl http://localhost/
cat /usr/share/nginx/html/index.html
echo 'Content inside NGINX File' > /usr/share/nginx/html/index.html

# To quit the shell in the container and go back to host machine
exit

$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
testapp-pod   1/1     Running   0           6m10s

# Running individual commands in a container
kubectl exec testapp-pod -- printenv
kubectl exec testapp-pod -- ls /

# Print the logs for a container in a pod
kubectl logs -f testapp-pod

# Delete all Evicted Pods
kubectl get pod | grep Evicted | awk '{print $1}' | xargs kubectl delete pod

```

ReplicaSets

- **ReplicaSet** is a way to setup replication of pods.
- ReplicaSet Definition File

replicaset-definition.yaml

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: testapp-replicaset
  labels:
    app: testapp
    type: front-end
spec:
  replicas: 3
  selector:
    matchLabels:
      type: front-end
  template:
    metadata:
      name: testapp-pod
      labels:

```

```

    app: testapp
    type: front-end
spec:
  containers:
  - name: nginx-container
    image: nginx

```

--

- Execute below commands:

```

# To apply the replicaset
kubectl apply -f replicaset-definition.yaml
$ kubectl get pods
# OUTPUT
NAME                                READY   STATUS    RESTARTS   AGE
testapp-replicaset-5bfm1            1/1     Running   0           14s
testapp-replicaset-9m58f            1/1     Running   0           14s
testapp-replicaset-g962z            1/1     Running   0           14s
$ kubectl get replicaset
$ kubectl get rs
# OUTPUT
NAME                DESIRED   CURRENT   READY   AGE
testapp-replicaset  3         3         3       17s

$ kubectl delete pod testapp-replicaset-5bfm1
pod "testapp-replicaset-5bfm1" deleted

$ kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
testapp-replicaset  3         3         3       116s

$ kubectl describe rs testapp-replicaset

$ kubectl get pods
# OUTPUT
NAME                                READY   STATUS    RESTARTS   AGE
testapp-replicaset-57rmw            1/1     Running   0           26s
testapp-replicaset-9m58f            1/1     Running   0           118s
testapp-replicaset-g962z            1/1     Running   0           118s

$ kubectl describe pod testapp-replicaset-57rmw

```

--

Scaling Replicaset

- There are multiple ways to scale replicaset:
 - 1. Update the number of replicas in the **replicaset-definition.yaml** definition file as **replicas: 5** and Re-apply the same file.

```
kubectl apply -f replicaset-definition.yaml
```

- 2. Use **kubectl scale** command with manifest yaml file.

```
kubectl scale --replicas=2 -f replicaset-definition.yaml
```

- 3. Use **kubectl scale** command with type and name.

```
kubectl scale --replicas=3 replicaset testapp-replicaset
```

Deployments

- Create Deployment with Imperative Command.

```
$ kubectl get deployments
```

```
$ kubectl create deployment hello-node --image=k8s.gcr.io/echoserver:1.4
```

```
$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-node	1/1	1	1	48s

```
$ kubectl describe deployment hello-node
```

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
hello-node-7567d9fdc9	1	1	1	50s

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-node-7567d9fdc9-cw8n5	1/1	Running	0	55s

```
$ kubectl describe pod hello-node-7567d9fdc9-cw8n5
```

Name: hello-node-7567d9fdc9-cw8n5

Namespace: default

Priority: 0

Node: minikube/10.0.0.7

Start Time: Tue, 26 Jul 2022 02:35:26 +0000

Labels: app=hello-node

pod-template-hash=7567d9fdc9

Annotations: <none>

Status: Running

IP: 172.18.0.6

IPs:

IP: 172.18.0.6

Controlled By: ReplicaSet/hello-node-7567d9fdc9


```
Containers:
  echoserver:
    Container ID:
docker://837d38e387b4d44eaa04b87ca00344a07e3170e4aa0f2f9df97edc22a0a5cfed
    Image:          k8s.gcr.io/echoserver:1.4
    Image ID:       docker-
pullable://gcr.io/google_containers/echoserver@sha256:5d99aa1120524c801bc8c1a7077e8f5ec122ba16b6dda1a5d3826057f67b9bcb
    Port:          <none>
    Host Port:     <none>
    State:         Running
      Started:     Tue, 26 Jul 2022 02:35:28 +0000
    Ready:         True
    Restart Count:  0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-4vfrr (ro)
Conditions:
  Type            Status
  Initialized     True
  Ready           True
  ContainersReady True
  PodScheduled    True
Volumes:
  default-token-4vfrr:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-4vfrr
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
Normal   Scheduled   2m3s  default-scheduler  Successfully assigned default/hello-node-7567d9fdc9-cw8n5 to minikube
Normal   Pulling     2m2s  kubelet         Pulling image
"k8s.gcr.io/echoserver:1.4"
Normal   Pulled      2m1s  kubelet         Successfully pulled image
"k8s.gcr.io/echoserver:1.4" in 487.451376ms
Normal   Created     2m1s  kubelet         Created container echoserver
Normal   Started     2m1s  kubelet         Started container echoserver

kubectl explain deployment
```

- Create deployment with **Declarative** Command, Download or use below command to apply the Deployment

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml
```

nginx-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

- **kind** - What kind of object you want to create.
- **metadata** - Data that helps uniquely identify the object, including a **name** string, **label**.
- **spec.replicas** — Tells Kubernetes how many pods to create during a deployment. Modifying this field is an easy way to scale a containerized application.
- **spec.selector** — An optional object that tells the Kubernetes deployment controller to only target pods that match the specified labels.
- **spec.template.metadata.labels** — Adds labels to a deployment specification.

--


Execute Commands : <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/#scaling-a-deployment>

Namespaces

- A Kubernetes namespace helps separate a cluster into logical units.
- In a new cluster, Kubernetes automatically creates the following namespaces: **default** (for user workloads) and for the Kubernetes control plane: **kube-system**.

- Kubernetes also allows admins to manually create custom namespaces.

Namespace - Isolation



•

--

Kubernetes Namespaces Concepts

- There are two types of Kubernetes namespaces: Kubernetes system namespaces and custom namespaces.
- **default** - a default space for objects that do not have a specified namespace.
- **kube-system** - a default space for Kubernetes system objects, such as kube-dns and kube-proxy, and add-ons providing cluster-level features, such as web UI dashboards, ingresses, and cluster-level logging.

--

- Use this command to list all the available namespaces in your environment.

```
kubectl get namespaces
```

```
kubectl get pods -n default
```

```
kubectl get pods -n kube-system
```

```
kubectl create namespace development
```

```
kubectl get namespaces
```

```
kubectl get pods -n development
```

```
# Launch a new pod in this newly created namespace
```

```
kubectl run new-nginx-pod --image nginx:1.16.1 -n development
```

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml -n
development

kubectl get deployments -n development
kubectl get pods -n development
kubectl get pods -n default
kubectl get pods

# Delete all pods in namespace
kubectl delete pods --all -n development

# To delete a namespace.
kubectl delete namespace development
```

Note: Kubernetes will always list the pods from the **default** namespace. you need to specify the namespace name to display the objects in it.

Kubernetes Service

- In the Kubernetes world, the pods, where the application lives, are temporary and get a new IP address every time they are launched.
- The pods are usually dynamically destroyed and recreated with each deployment.
- In the absence of the Kubernetes service, we would have to track the IP addresses of all active pods.
- The **Kubernetes service** creates an abstraction that maps to one or more pods.
- This abstraction allows other applications to reach the service by simply referring to the service name.
- It means that other applications no longer need to know the IP addresses assigned to the pods.
- External applications and end-users can also access the services assuming that they are exposed publicly to the internet.
- Kubernetes Services enables communication between various components within and outside of the application Container Pods
- A Service enables network access to a set of Pods in Kubernetes.

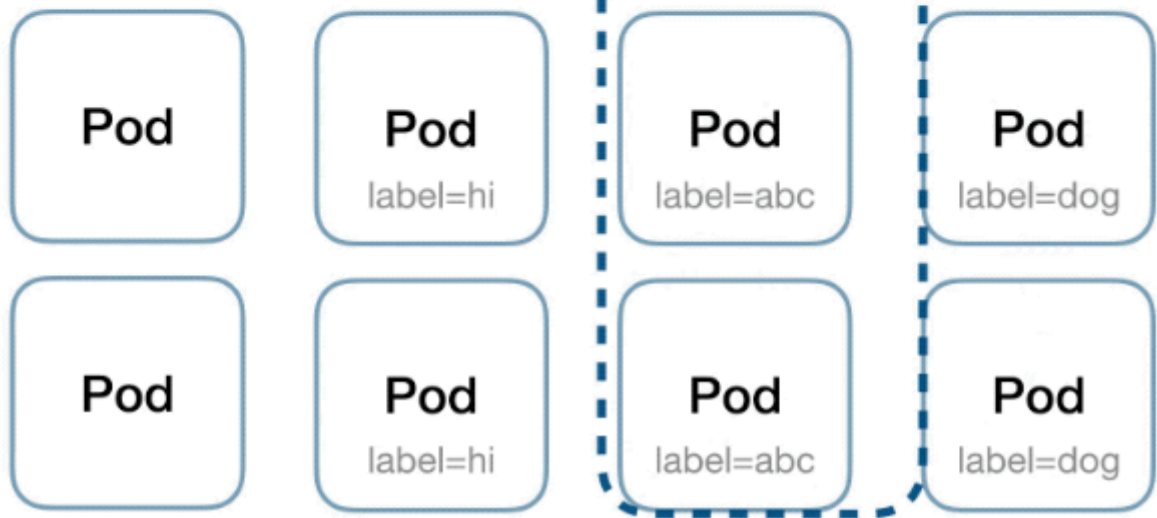
--

- **Services** match a set of Pods using **labels** and **selectors**.
- Labels are key/value pairs attached to objects and can be used in any number of ways:
 - Designate objects for **development, test, and production**
 - Classify an object using **tags**
- When a network request is made to the **service**, it selects all **Pods** in the cluster matching the **service's selector**, chooses one of them, and forwards the network request to it.

--

Service

Select a pod where
label = abc



Network
Request

- There are 3 types of service types in kubernetes.

Services Types



NodePort

- Exposes the Service on each Node's IP at a static port (the NodePort).
- If it is possible to contact the NodePort Service, from outside the cluster, by requesting **NodeIP:NodePort**
- A **ClusterIP** Service, to which the NodePort Service routes, is automatically created.
- In NodePort, the service makes an internal POD accessible on a PORT on the NODE.

nginx-deployment-definition.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  # Create 3 replicas of the Pods.
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
```

```

- name: nginx
  image: nginx:1.14.2
  ports:
    - containerPort: 80

```

service-definition.yaml

```

apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  # Expose the service on a static port on each node
  # so that we can access the service from outside the cluster
  type: NodePort

  # When the node receives a request on the static port (30163)
  # "select pods with the label 'app' set to 'echo-hostname'"
  # and forward the request to one of them
  selector:
    app: nginx

  ports:
    # Three types of ports for a service
    # nodePort - a static port assigned on each the node
    # port - port exposed internally in the cluster
    # targetPort - the container port to send requests to
    - nodePort: 30163
      port: 8080
      targetPort: 80

```

--

- Use **kubectl apply** to apply both above YAML Manifest Files

```

kubectl apply -f nginx-deployment-definition.yaml
kubectl get pods -o wide
kubectl apply -f service-definition.yaml
kubectl get svc
# OUTPUT
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes           ClusterIP     10.96.0.1     <none>         443/TCP          24m
nginx-service        NodePort      10.101.248.0  <none>         8080:30163/TCP   6m20s

# Access the Port 30163 in the View Port

$ kubectl describe service nginx-service
# OUTPUT
Name:                nginx-service

```

```
Namespace:          default
Labels:             <none>
Annotations:        <none>
Selector:           app=nginx
Type:               NodePort
IP Families:        <none>
IP:                 10.107.15.76
IPs:                10.107.15.76
Port:               <unset> 8080/TCP
TargetPort:         80/TCP
NodePort:           <unset> 30163/TCP
Endpoints:          10.5.1.6:80,10.5.1.7:80,10.5.1.8:80
Session Affinity:   None
External Traffic Policy: Cluster
Events:             <none>
```

The Endpoints in the above Service Description are the POD IPs.
Each time page is refreshed, the response of the WebPage is from any one of the POD that is running with Label "app: nginx"

```
$ kubectl describe endpoints nginx-service
```

```
# OUTPUT
```

```
Name:          nginx-web-service
Namespace:     default
Labels:        <none>
Annotations:   endpoints.kubernetes.io/last-change-trigger-time: 2023-04-
03T19:35:42Z
Subsets:
  Addresses:          10.5.1.6,10.5.1.7,10.5.1.8
  NotReadyAddresses:  <none>
  Ports:
    Name      Port  Protocol
    ----      -
    <unset>   80    TCP
```

```
Events:  <none>
```

Login inside one of the Pod to modify the /usr/share/nginx/html/index.html

```
kubectl get pods -o wide --show-labels
```

```
kubectl exec -it nginx-deployment-66b6c48dd5-kjlws /bin/bash
```

```
root@nginx-deployment-66b6c48dd5-2k7w8:/# cd /usr/share/nginx/html/
```

```
root@nginx-deployment-66b6c48dd5-2k7w8:/# ls
```

```
root@nginx-deployment-66b6c48dd5-2k7w8:/# echo "Modified File inside one POD with
Hostname as $HOSTNAME" > index.html
```

```
root@nginx-deployment-66b6c48dd5-2k7w8:/# exit
```

Access the Port 30163 in the View Port/Refresh the page multiple times.

```
curl localhost:30163
```

```
curl localhost:30163
```

```
curl localhost:30163
```

The Service will forward request to the Pod for which above Page is modified.

--

```
kind: Service
apiVersion: v1
```

```
metadata:
```

```
  name: hostname-service
```

```
spec:
```

```
  type: NodePort
```

```
  selector:
```

```
    app: echo-hostname
```

```
  ports:
```

```
  - nodePort: 30163
```

```
    port: 8080
```

```
    targetPort: 80
```

Make the service available to network requests from external clients

Forward requests to pods with label of this value

nodePort

access service via this external port number

port

port number exposed internally in cluster

targetPort

port that containers are listening on

--

ClusterIP

- By default, Kubernetes creates a **ClusterIP** type of service. We can build different kinds of services by having a **spec.type** property in the service YAML file.
- Exposes the Service on a cluster-internal IP.
- Choosing this value makes the Service only reachable from within the cluster. This is the default **ServiceType**.

pod-service-definition.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-web-deployment
  labels:
    app: nginx-web
spec:
  # Create 3 replicas of the Pods.
  replicas: 3
  selector:
    matchLabels:
      app: nginx-web
  template:
```

```

  metadata:
    labels:
      app: nginx-web
  spec:
    containers:
      - name: nginx-web
        image: nginx:1.14.2
        ports:
          - containerPort: 80

# ---

apiVersion: v1
kind: Service
metadata:
  name: nginx-web-service
spec:
  # Expose the service on a static port on each node
  # so that we can access the service from outside the cluster
  type: ClusterIP

  # When the node receives a request on the static port (30163)
  # "select pods with the label 'app' set to 'echo-hostname'"
  # and forward the request to one of them
  selector:
    app: nginx-web

  ports:
    # Three types of ports for a service
    # port - port exposed internally in the cluster
    # targetPort - the container port to send requests to
    - port: 8080
      targetPort: 80

```

```

kubectl apply -f pod-service-definition.yaml
# OUTPUT
deployment.apps/nginx-web-deployment unchanged
service/nginx-web-service configured

```

```

kubectl get pods -o wide
$ kubectl get pods -o wide

```

NAME		READY	STATUS	RESTARTS	AGE	IP		
NODE	NOMINATED NODE	READINESS GATES						
nginx-web-deployment-7949b98ff9-9hhzg	172.18.0.10	minikube	<none>	1/1	Running	0	9m13s	
nginx-web-deployment-7949b98ff9-ht64b	172.18.0.9	minikube	<none>	1/1	Running	0	9m13s	
nginx-web-deployment-7949b98ff9-s9qjh	172.18.0.11	minikube	<none>	1/1	Running	0	9m13s	

```

$ kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
54m				
nginx-web-service	ClusterIP	10.108.233.187	<none>	8080/TCP
8m58s				

```
$ kubectl describe service nginx-web-service
Name:          nginx-web-service
Namespace:     default
Labels:        <none>
Annotations:   <none>
Selector:      app=nginx-web
Type:          ClusterIP
IP Families:   <none>
IP:            10.108.233.187
IPs:           10.108.233.187
Port:          <unset> 8080/TCP
TargetPort:    80/TCP
Endpoints:     172.18.0.10:80,172.18.0.11:80,172.18.0.9:80
Session Affinity: None
Events:        <none>

curl 10.108.233.187:8080
curl 10.108.233.187:8080
```

--

LoadBalancer

- This service type creates load balancers in various Cloud providers like AWS, GCP, Azure, etc., to expose our application to the Internet.
- The Cloud provider will provide a mechanism for routing the traffic to the services. The most common example usage of this type is for a website or a web app.
- On cloud providers which support external load balancers, setting the type field to LoadBalancer provisions a load balancer for your Service.

pod-service-definition.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-web-service
spec:
  # Expose the service on a static port on each node
  # so that we can access the service from outside the cluster
  type: LoadBalancer

  # When the node receives a request on the static port (30163)
  # "select pods with the label 'app' set to 'echo-hostname'"
```

```
# and forward the request to one of them
selector:
  app: nginx-web

ports:
  # Three types of ports for a service
  # port - port exposed internally in the cluster
  # targetPort - the container port to send requests to
  - protocol: TCP
    port: 8080
    targetPort: 80
```

- Traffic from the external load balancer is directed at the backend Pods. The cloud provider decides how it is load balanced.

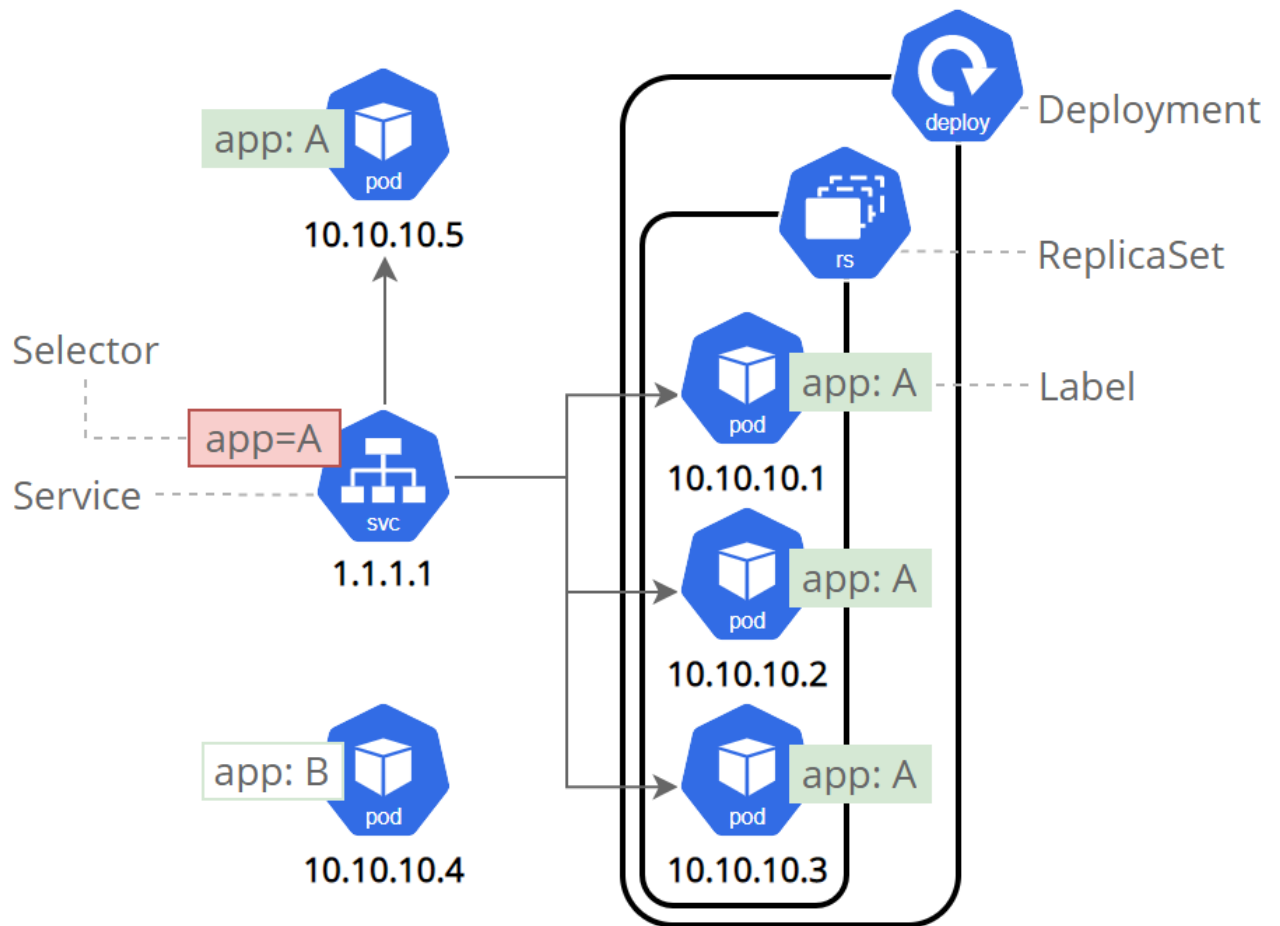
--

Kubernetes Service vs Deployment

- A deployment is responsible for keeping a set of pods running.
- A service is responsible for enabling network access to a set of pods.
- A deployment without a service can be created to keep a set of identical pods running in the Kubernetes cluster.
- The deployment could be scaled up and down and pods could be replicated.
- Each pod could be accessed individually via direct network requests (rather than abstracting them behind a service), but keeping track of this for a lot of pods is difficult.
- Services and Deployments are different, but they work together to route network requests across pods.

--

- Kubernetes Objects have labels and selectors as below:



•

Reference

Kubernetes Cluster Setup

- [Minikube](#) – Development and Learning
- [Kops](#) – Learning, Development, Production
- [Kubeadm](#) – Learning, Development, Production
- [Docker for Mac](#) - Learning, Development
- [Kubernetes IN Docker](#) - Learning, Development