

Table of Contents

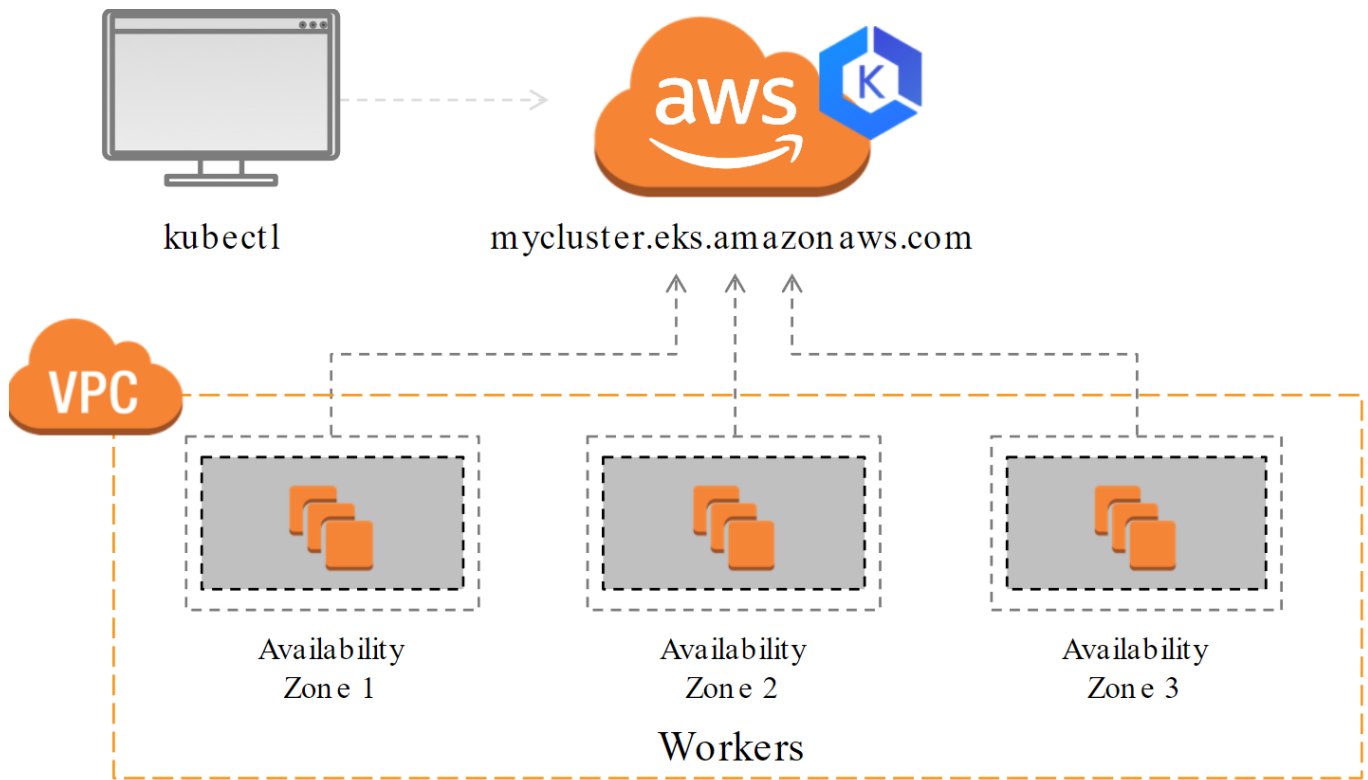
- [Table of Contents](#)
 - [What is EKS](#)
 - [EKS ARCHITECTURE FOR CONTROL PLANE AND WORKER NODE COMMUNICATION](#)
 - [Steps to create EKS](#)
 - [EKSCTL CLI](#)
 - [Verify Cluster, NodeGroup in EKS Management Console](#)
 - [List Worker Nodes](#)
 - [kubectl commands](#)
 - [Kubernetes Application Deployment](#)
 - [Docker Image Creation and Push to Registry](#)
 - [EKS Cluster Pricing](#)
 - [EKS Cluster Pricing](#)
 - [EKS Worker Nodes Pricing-EC2](#)
 - [EKS Case Study:](#)
 - [Reference](#)
 - [Reference:](#)
-

What is EKS

- AWS EKS is a **Managed Kubernetes Service** from Amazon, which means AWS manages the Master Nodes.
- All the necessary applications/services are already pre-installed like the container runtime or master processes and in addition it also takes care of scaling and backups.
- You only create the Worker Nodes which are EC2 instances.

--

EKS ARCHITECTURE FOR CONTROL PLANE AND WORKER NODE COMMUNICATION



--

Steps to create EKS

- To create a K8s cluster in EKS you need to do following steps:
- Create a VPC
- Create an IAM role with Security Group (or in other words: create AWS user with list of permissions)
- Create Cluster Control Plane - Master Nodes
 - choose basic information like cluster name and k8s version
 - choose region and VPC for your cluster
 - set security
- Create Worker Nodes and connect to cluster
 - The Worker Nodes are some EC2 instances with CPU and storage resources.
 - Create as a Node Group
 - Choose cluster it will attach to
 - Define Security Group, select instance type etc.

With **NodeGroup** you have autoscaling, depending on how much load the cluster has new Worker Nodes will automatically added or removed in the cluster.

--

EKSCTL CLI

- Execute below commands in either one of the below options:
 - AWS CloudShell (Uses AWS Console Login IAM Credentials)
 - Local Linux Machine (Uses IAM User Credentials configured)
 - Amazon EC2 Linux Machine (Uses AWS IAM Role Permissions)

- Download the `eksctl` on your machine linux cli

```
curl --silent --location
"https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -
s)_amd64.tar.gz" | tar xz -C /tmp
sudo mv /tmp/eksctl /usr/local/bin
eksctl version
```

- The IAM account used for EKS cluster creation should have these minimal access levels.

AWS Service	Access Level
CloudFormation	Full Access
EC2	Full: Tagging Limited: List, Read, Write
EC2 Auto Scaling	Limited: List, Write
EKS	Full Access
IAM	Limited: List, Read, Write, Permissions Management
Systems Manager	Limited: List, Read

--

- Command to create EKS Cluster

```
eksctl create cluster \
  --name test-eks-cluster \
  --version 1.25 \
  --region ap-south-1 \
  --nodegroup-name eks-worker-nodegroup \
  --node-type t4g.small \
  --nodes 2

#####OUTPUT#####
2022-01-26 01:53:53 [i] eksctl version 0.80.0
2022-01-26 01:53:53 [i] using region ap-south-1
2022-01-26 01:53:54 [i] setting availability zones to [ap-south-1b ap-south-1a
ap-south-1c]
2022-01-26 01:53:54 [i] subnets for ap-south-1b - public:192.168.0.0/19
private:192.168.96.0/19
2022-01-26 01:53:54 [i] subnets for ap-south-1a - public:192.168.32.0/19
private:192.168.128.0/19
2022-01-26 01:53:54 [i] subnets for ap-south-1c - public:192.168.64.0/19
private:192.168.160.0/19
2022-01-26 01:53:54 [i] nodegroup "ng-2564739d" will use "" [AmazonLinux2/1.21]
2022-01-26 01:53:54 [i] using Kubernetes version 1.21
2022-01-26 01:53:54 [i] creating EKS cluster "test-eks-cluster" in "ap-south-1"
region with managed nodes
2022-01-26 01:53:54 [i] will create 2 separate CloudFormation stacks for cluster
```

```

itself and the initial managed nodegroup
2022-01-26 01:53:54 [i] if you encounter any issues, check CloudFormation console
or try 'eksctl utils describe-stacks --region=ap-south-1 --cluster=test-eks-
cluster'
2022-01-26 01:53:54 [i] Kubernetes API endpoint access will use default of
{publicAccess=true, privateAccess=false} for cluster "test-eks-cluster" in "ap-
south-1"
2022-01-26 01:53:54 [i] CloudWatch logging will not be enabled for cluster "test-
eks-cluster" in "ap-south-1"
2022-01-26 01:53:54 [i] you can enable it with 'eksctl utils update-cluster-
logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=ap-south-
1 --cluster=test-eks-cluster'
2022-01-26 01:53:54 [i]
2 sequential tasks: { create cluster control plane "test-eks-cluster",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create managed nodegroup "ng-2564739d",
  }
}
2022-01-26 01:53:54 [i] building cluster stack "eksctl-test-eks-cluster-cluster"
2022-01-26 01:53:55 [i] deploying stack "eksctl-test-eks-cluster-cluster"
2022-01-26 01:54:25 [i] waiting for CloudFormation stack "eksctl-test-eks-
cluster-cluster"

2022-10-19 02:37:29 [i] waiting for the control plane to become ready
2022-10-19 02:37:30 [✓] saved kubeconfig as "/home/cloudshell-
user/.kube/config"
2022-10-19 02:37:30 [i] no tasks
2022-10-19 02:37:30 [✓] all EKS cluster resources for "test-eks-cluster" have
been created
2022-10-19 02:37:31 [i] nodegroup "eks-worker-nodegroup" has 2 node(s)
2022-10-19 02:37:31 [i] node "ip-192-168-31-69.ap-south-1.compute.internal" is
ready
2022-10-19 02:37:31 [i] node "ip-192-168-62-91.ap-south-1.compute.internal" is
ready
2022-10-19 02:37:31 [i] waiting for at least 2 node(s) to become ready in "eks-
worker-nodegroup"
2022-10-19 02:37:31 [i] nodegroup "eks-worker-nodegroup" has 2 node(s)
2022-10-19 02:37:31 [i] node "ip-192-168-31-69.ap-south-1.compute.internal" is
ready
2022-10-19 02:37:31 [i] node "ip-192-168-62-91.ap-south-1.compute.internal" is
ready
2022-10-19 02:37:34 [i] kubectl command should work with "/home/cloudshell-
user/.kube/config", try 'kubectl get nodes'
2022-10-19 02:37:34 [✓] EKS cluster "test-eks-cluster" in "ap-south-1" region
is ready

```

```
#####OUTPUT#####
```

--

- Install **kubectl** linux utility on your local linux client.

```
curl -o kubectl https://s3.us-west-2.amazonaws.com/amazon-eks/1.21.5/2022-01-21/bin/linux/amd64/kubectl
chmod +x ./kubectl
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export
PATH=$PATH:$HOME/bin
echo 'export PATH=$PATH:$HOME/bin' >> ~/.bashrc
```

```
# View the Local kubeconfig file
kubectl config view
# This command shows content inside the ~/.kube/config file
#####OUTPUT#####
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+OMITTED
    server: https://00D6914D1CB90D39116178B3A9A2ECA6.y14.ap-south-1.eks.amazonaws.com
    name: test-eks-cluster.ap-south-1.eksctl.io
contexts:
- context:
    cluster: test-eks-cluster.ap-south-1.eksctl.io
    user: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
    name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
current-context: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
kind: Config
preferences: {}
users:
- name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - eks
        - get-token
        - --cluster-name
        - test-eks-cluster
        - --region
        - ap-south-1
      command: aws
      env:
        - name: AWS_STS_REGIONAL_ENDPOINTS
          value: regional
      provideClusterInfo: false
#####OUTPUT#####
```

--

When you create an Amazon EKS cluster, the AWS Identity and Access Management (IAM) entity user or role that creates the cluster, is automatically granted `system:masters` permissions in the cluster's

role-based access control (RBAC) configuration in the Amazon EKS control plane. This IAM entity doesn't appear in any visible configuration, so make sure to keep track of which IAM entity originally created the cluster.

- Validate the CloudFormation Template Resources created with above command.
 - Stack Name : **eksctl-test-eks-cluster-cluster**
 - Stack Name : **eksctl-test-eks-cluster-nodegroup-eks-worker-nodegroup**

--

Verify Cluster, NodeGroup in EKS Management Console

- Go to Services -> **Elastic Kubernetes Service**

List Worker Nodes

```
# List EKS clusters
eksctl get cluster

# List NodeGroups in a cluster
eksctl get nodegroup --cluster=<clusterName>

# List Nodes in current kubernetes cluster
kubectl get nodes
[cloudshell-user@ip-10-0-173-52 ~]$ kubectl get nodes
Kubeconfig user entry is using deprecated API version
client.authentication.k8s.io/v1alpha1. Run 'aws eks update-kubeconfig' to update.
NAME                                                    STATUS    ROLES    AGE     VERSION
ip-192-168-4-237.ap-south-1.compute.internal            Ready     <none>    7m58s   v1.21.12-eks-5308cf7
ip-192-168-85-16.ap-south-1.compute.internal            Ready     <none>    8m5s    v1.21.12-eks-5308cf7

# Our kubectl context should be automatically changed to new cluster
kubectl config view --minify
```

--

kubectl commands

```
kubectl get pods -n kube-system
kubectl get pods -n kube-system -o wide
```

- Create the Deployment by running the following commands:
 - Reference : [run-stateless-application-deployment](#)

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml
wget https://k8s.io/examples/controllers/nginx-deployment.yaml
```

```
[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get events
```

LAST SEEN	TYPE	REASON	OBJECT
MESSAGE			
2m48s	Normal	Scheduled	pod/nginx-deployment-66b6c48dd5-qkst9
Successfully assigned default/nginx-deployment-66b6c48dd5-qkst9 to ip-192-168-21-139.ap-south-1.compute.internal			
2m47s	Normal	Pulling	pod/nginx-deployment-66b6c48dd5-qkst9
Pulling image "nginx:1.14.2"			
2m38s	Normal	Pulled	pod/nginx-deployment-66b6c48dd5-qkst9
Successfully pulled image "nginx:1.14.2" in 9.12123707s			
2m37s	Normal	Created	pod/nginx-deployment-66b6c48dd5-qkst9
Created container nginx			
2m37s	Normal	Started	pod/nginx-deployment-66b6c48dd5-qkst9
Started container nginx			
37s	Warning	FailedScheduling	pod/nginx-deployment-66b6c48dd5-x92pj
0/2 nodes are available: 2 Too many pods.			
2m48s	Normal	Scheduled	pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully assigned default/nginx-deployment-66b6c48dd5-zb9bb to ip-192-168-46-56.ap-south-1.compute.internal			
2m47s	Normal	Pulling	pod/nginx-deployment-66b6c48dd5-zb9bb
Pulling image "nginx:1.14.2"			
2m37s	Normal	Pulled	pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully pulled image "nginx:1.14.2" in 9.470360458s			
2m37s	Normal	Created	pod/nginx-deployment-66b6c48dd5-zb9bb
Created container nginx			
2m37s	Normal	Started	pod/nginx-deployment-66b6c48dd5-zb9bb
Started container nginx			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-zb9bb			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-qkst9			
2m48s	Normal	SuccessfulCreate	replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-x92pj			
2m48s	Normal	ScalingReplicaSet	deployment/nginx-deployment
Scaled up replica set nginx-deployment-66b6c48dd5 to 3			

```
kubectl describe pods -n default > pods.txt
cat pods.txt
```

```
Name:          nginx-deployment-66b6c48dd5-qkst9
Namespace:     default
Priority:       0
Node:          ip-192-168-21-139.ap-south-1.compute.internal/192.168.21.139
Start Time:    Wed, 02 Feb 2022 19:29:23 +0000
Labels:        app=nginx
               pod-template-hash=66b6c48dd5
Annotations:   kubernetes.io/psp: eks.privileged
Status:        Running
IP:            192.168.15.247
IPs:
```

IP: 192.168.15.247
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5

Name: nginx-deployment-66b6c48dd5-x92pj
Namespace: default
Priority: 0
Node: <none>
Labels: app=nginx
pod-template-hash=66b6c48dd5
Annotations: kubernetes.io/psp: eks.privileged
Status: Pending

IP: <none>
IPs: <none>
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5

Name: nginx-deployment-66b6c48dd5-zb9bb
Namespace: default
Priority: 0
Node: ip-192-168-46-56.ap-south-1.compute.internal/192.168.46.56
Start Time: Wed, 02 Feb 2022 19:29:23 +0000
Labels: app=nginx
pod-template-hash=66b6c48dd5
Annotations: kubernetes.io/psp: eks.privileged
Status: Running
IP: 192.168.51.252
IPs: 192.168.51.252
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5

kubectl get pods -o wide

- The above command shows pods details where the pod gets an IP from ENI in the Subnet.

Network interfaces (6) Info

Filter network interfaces

search: eks Clear filters

Subnet ID	VPC ID	Availability Zone	Interface Type	Description	Instance ID	Status	Public IPv4...	Primary private IPv4...	Secondary private IPv4 add...
subnet-067941fe3a2fce5ab	vpc-02b76a6ce27d533a6	ap-south-1c	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.165.131	-
subnet-00fc8a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	aws-K8S-I-08266ca7c...	i-08266ca7c22b644a5	In-use	-	192.168.30.4	192.168.15.247
subnet-00fc8a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	-	i-08266ca7c22b644a5	In-use	65.0.129.57	192.168.21.139	192.168.7.193
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	aws-K8S-I-0c544c004...	i-0c544c004ed1ea632	In-use	-	192.168.49.221	192.168.51.252
subnet-004a60aba0384c11a	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.150.81	-
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	-	i-0c544c004ed1ea632	In-use	3.111.34.130	192.168.46.56	192.168.46.207

- AWS EKS supports native VPC networking with the Amazon VPC Container Network Interface (CNI) plugin for Kubernetes.
- Using this plugin allows Kubernetes Pods to have the same IP address inside the pod as they do on the VPC network.
- For more information, see amazon-vpc-cni-k8s and Proposal: CNI plugin for Kubernetes networking over AWS VPC on GitHub.

- The Amazon VPC CNI plugin is fully supported for use on Amazon EKS and self-managed Kubernetes clusters on AWS.
- Refer the ENI Limit as per Instance Type :
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html#AvailableIpPerENI>
 - The formula for defining the maximum number of Pods per EC2 Node instance is as follows:
 - $N * (M-1) + 2$
 - **N** is the number of Elastic Network Interfaces (ENI) of the instance type
 - **M** is the number of IP addresses per ENI.
 - For e.g. **t2.micro** instance, this calculation is $2 * (2-1) + 2 = 4$ Pods

--

- Use below command to get similar details:

```
aws ec2 describe-instance-types --filters "Name=instance-type,Values=t4g.*" --
query "InstanceTypes[].[Type: InstanceType, MaxENI:
NetworkInfo.MaximumNetworkInterfaces, IPv4addr:
NetworkInfo.Ipv4AddressesPerInterface]" --output table
```

DescribeInstanceTypes			
IPv4addr	MaxENI	Type	
2	2	t4g.nano	
6	3	t4g.medium	
12	3	t4g.large	
15	4	t4g.2xlarge	
4	3	t4g.small	
2	2	t4g.micro	
15	4	t4g.xlarge	

```
kubectl get pods --all-namespaces -o wide
kubectl get pods --all-namespaces -o wide | grep -i running
[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get pods --all-namespaces -o wide |
grep -i running
default      nginx-deployment-66b6c48dd5-qkst9  1/1      Running   0   38m
192.168.15.247 ip-192-168-21-139.ap-south-1.compute.internal
default      nginx-deployment-66b6c48dd5-zb9bb  1/1      Running   0   38m
192.168.51.252 ip-192-168-46-56.ap-south-1.compute.internal
kube-system  aws-node-jls58                      1/1      Running   0   125m
192.168.46.56 ip-192-168-46-56.ap-south-1.compute.internal
kube-system  aws-node-tmw12                      1/1      Running   0   125m
192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system  coredns-7f95bc96cc-dl6dn           1/1      Running   0   135m
192.168.46.207 ip-192-168-46-56.ap-south-1.compute.internal
kube-system  coredns-7f95bc96cc-fj56b           1/1      Running   0   135m
192.168.7.193 ip-192-168-21-139.ap-south-1.compute.internal
```

```
kube-system    kube-proxy-dqlph          1/1    Running   0    125m
192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system    kube-proxy-xlpfk          1/1    Running   0    125m
192.168.46.56 ip-192-168-46-56.ap-south-1.compute.internal
[cloudshell-user@ip-10-0-50-112 ~]$
```

--

Kubernetes Application Deployment

Docker Image Creation and Push to Registry

- Below Steps can be performed on any Docker Host or CI Job.
- Clone Repo: <https://github.com/devops-practicals/notes-app-cicd.git>

```
cd notes-app-cicd

sudo docker login
# Enter DockerHub Username and Password

# Build Docker Image

sudo docker build -t notes-app .

# Tag Docker Image
sudo docker tag notes-app:latest <Dockerhub_Username>/notes-app:latest

# Push the Docker image
sudo docker push <Dockerhub_Username>/notes-app:latest
```

- Once Image is available in the Image Registry i.e DockerHub/ECR, specify this image Name:Tag in the Kubernetes Deployment YAML Manifest Files.

```
# Validate the image in deployment.yaml file
apiVersion: apps/v1
kind: Deployment
metadata:
  name: notes-app-deployment
  labels:
    app: notes-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: notes-app
  template:
    metadata:
      labels:
        app: notes-app
```

```

spec:
  containers:
  - name: notes-app-deployment
    image: <Dockerhub_Username>/notes-app:latest
    resources:
      requests:
        cpu: "100m"
    imagePullPolicy: IfNotPresent
    ports:
      - containerPort: 3000

# Validate the service.yaml file
apiVersion: v1
# Indicates this as a service
kind: Service
metadata:
  # Service name
  name: notes-app-deployment
spec:
  selector:
    # Selector for Pods
    app: notes-app
  ports:
    # Port Map
    - port: 80
      targetPort: 3000
      protocol: TCP
  type: LoadBalancer

# Apply the manifest files with the following commands. Starting with deployment
# and then service yaml file.
sudo kubectl apply -f deployment.yaml

sudo kubectl apply -f service.yaml

sudo kubectl get pods
#
NAME                                READY   STATUS    RESTARTS   AGE
notes-app-deployment-568df5649-5vdjc 1/1     Running   0           16s
notes-app-deployment-568df5649-g286v 1/1     Running   0           16s

sudo kubectl get svc
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)
AGE
kubernetes                         ClusterIP            10.96.0.1        <none>            443/TCP
14m
notes-app-deployment               LoadBalancer        10.111.50.144    <pending>         80:31263/TCP
5s

# Access the application UI on Host Public IP with port as 31263
# http://3.110.78.93:31263/

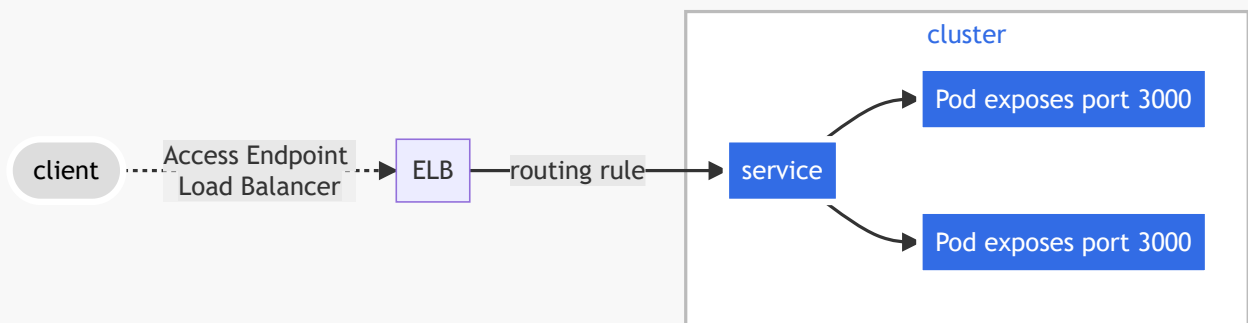
# If you executed this Kubernetes deployment and service on EKS Cluster, a ELB
# Endpoint will be created.

```

```
# Access the application on the ELB Endpoint.
[cloudshell-user@ip-10-2-86-100 ~]$ kubectl describe svc notes-app-deployment
Name:                notes-app-deployment
Namespace:           default
Labels:              <none>
Annotations:         <none>
Selector:            app=notes-app
Type:                LoadBalancer
IP Family Policy:    SingleStack
IP Families:         IPv4
IP:                  10.100.113.118
IPs:                 10.100.113.118
LoadBalancer Ingress: a236feab6f55f47d99c3fb72709686f9-1489323979.ap-south-1.elb.amazonaws.com
Port:                <unset> 80/TCP
TargetPort:          3000/TCP
NodePort:            <unset> 31263/TCP
Endpoints:           192.168.28.162:3000,192.168.43.57:3000
Session Affinity:    None
External Traffic Policy: Cluster
Events:
  Type    Reason              Age   From                    Message
  ----    -
  Normal  EnsuringLoadBalancer 43m   service-controller     Ensuring load balancer
  Normal  EnsuredLoadBalancer  43m   service-controller     Ensured load balancer
```

If you delete the kubernetes service object, the ELB endpoint is deleted.

```
kubectl delete svc notes-app-deployment
```



Once Testing of EKS Cluster Creation is done, make sure to Delete the cluster when not in use.

```
eksctl delete cluster --name test-eks-cluster
```

EKS Cluster Pricing

- EKS is not free (Unlike other AWS Services), In short, no free-tier for EKS.

EKS Cluster Pricing

- Below is the EKS Cluster Pricing:
 - \$0.10/hour
 - \$2.4/day
 - \$72/month

EKS Worker Nodes Pricing-EC2

- You pay for AWS resources (e.g. EC2 instances or EBS volumes)
- For **t3.medium** in N.Virginia
 - \$0.0416 per Hour
 - Per Day: \$0.9984 - Approximately \$1
 - Per Month: \$30 per 1 t3.medium Node.
- Reference: [EC2 On-Demand Pricing](#)

--

EKS Case Study:

- [meity-gov-india-case-study](#)
- [Architectural lessons from CoWIN platform](#)
- [cowin-public-v2](#)

--

Reference

- [EKS Cost Optimization](#)
- [Kubernetes instance calculator](#)
- [Resource Limits and Requests](#)

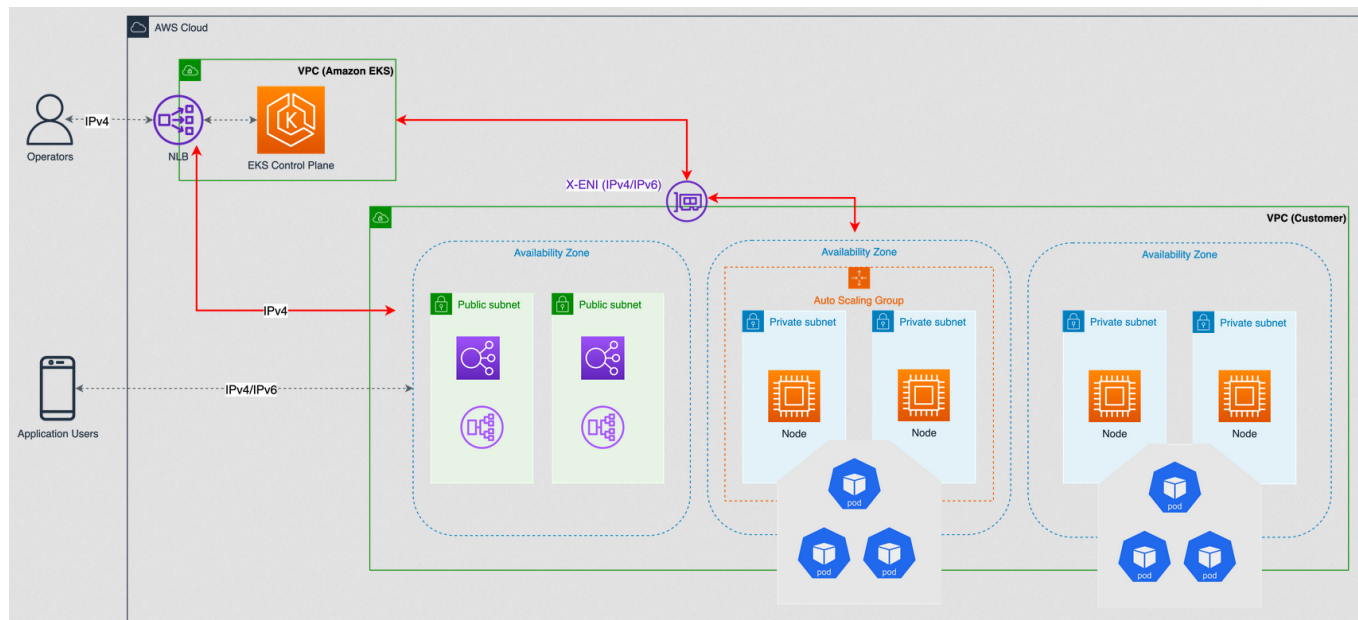
--

Reference:

- [AWS Graviton2 t4g.small free for all AWS Accounts until December 31, 2023](#)
 - Until December 31, 2023, all AWS customers will be enrolled automatically in the t4g.small Free Trial as detailed in the AWS Free Tier.
 - During the free-trial period, customers who run a t4g.small instance will automatically get 750 free hours per month deducted from their bill during each month.

--

- [aws-eks-best-practices](#)



- eks-cluster-connection