

# Table of Contents

---

- [Table of Contents](#)
  - [What is EKS](#)
  - [EKS Cluster Creation using Terraform](#)
    - [Verify Cluster, NodeGroup in EKS Management Console](#)
    - [List Worker Nodes](#)
    - [kubectl commands](#)
  - [Reference](#)
  - [EKS Cluster Pricing](#)
    - [EKS Cluster Pricing](#)
    - [EKS Worker Nodes Pricing-EC2](#)
  - [EKS Case Study:](#)
  - [Reference:](#)

## What is EKS

- AWS EKS is a **Managed Kubernetes Service** from Amazon, which means AWS manages the Master Nodes.
- All the necessary applications/services are already pre-installed like the container runtime or master processes and in addition it also takes care of scaling and backups.
- You only create the Worker Nodes which are EC2 instances.

## EKS Cluster Creation using Terraform

- For creating the EKS cluster using Terraform, Check for Latest Kubernetes Version [here](#)
- Follow steps mentioned in this Document:  
<https://developer.hashicorp.com/terraform/tutorials/kubernetes/eks>
- Modify the **eks-cluster.tf** to include **instance\_types = ["t4g.small"]**
- Modify the **eks-cluster.tf** to only have one **NodeGroup**
- Execute:
  - terraform init
  - terraform plan
  - terraform apply
- Validate all resources created in AWS Region **us-east-2**.
- Make sure **teraform destroy** is executed to delete all infrastructure.
- <https://github.com/AdminTurnedDevOps/Kubernetes-Quickstart-Environments/tree/main/aws>
- Install **kubect1** linux utility on your local linux client.

```
curl -o kubectl https://s3.us-west-2.amazonaws.com/amazon-eks/1.21.5/2022-01-21/bin/linux/amd64/kubectl
chmod +x ./kubectl
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export
PATH=$PATH:$HOME/bin
echo 'export PATH=$PATH:$HOME/bin' >> ~/.bashrc
```

```
# View the Local kubeconfig file
kubectl config view
# This command shows content inside the ~/.kube/config file
#####OUTPUT#####
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+OMITTED
    server: https://00D6914D1CB90D39116178B3A9A2ECA6.y14.ap-south-1.eks.amazonaws.com
    name: test-eks-cluster.ap-south-1.eksctl.io
contexts:
- context:
    cluster: test-eks-cluster.ap-south-1.eksctl.io
    user: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
    name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
current-context: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
kind: Config
preferences: {}
users:
- name: ansible-user@test-eks-cluster.ap-south-1.eksctl.io
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      args:
        - eks
        - get-token
        - --cluster-name
        - test-eks-cluster
        - --region
        - ap-south-1
      command: aws
      env:
        - name: AWS_STS_REGIONAL_ENDPOINTS
          value: regional
      provideClusterInfo: false
#####OUTPUT#####
```

When you create an Amazon EKS cluster, the AWS Identity and Access Management (IAM) entity user or role that creates the cluster, is automatically granted **system:masters** permissions in the cluster's **role-based access control** (RBAC) configuration in the Amazon EKS control plane. This IAM entity

doesn't appear in any visible configuration, so make sure to keep track of which IAM entity originally created the cluster.

- Validate the CloudFormation Template Resources created with above command.
  - Stack Name : **eksctl-test-eks-cluster-cluster**
  - Stack Name : **eksctl-test-eks-cluster-nodgroup-eks-worker-nodgroup**

## Verify Cluster, NodeGroup in EKS Management Console

- Go to Services -> **Elastic Kubernetes Service**

## List Worker Nodes

```
# List EKS clusters
eksctl get cluster

# List NodeGroups in a cluster
eksctl get nodegroup --cluster=<clusterName>

# List Nodes in current kubernetes cluster
kubectl get nodes
[cloudshell-user@ip-10-0-173-52 ~]$ kubectl get nodes
Kubeconfig user entry is using deprecated API version
client.authentication.k8s.io/v1alpha1. Run 'aws eks update-kubeconfig --region
region-code --name my-cluster' to update.
NAME                                STATUS    ROLES    AGE    VERSION
ip-192-168-4-237.ap-south-1.compute.internal Ready    <none>   7m58s  v1.21.12-
eks-5308cf7
ip-192-168-85-16.ap-south-1.compute.internal Ready    <none>   8m5s   v1.21.12-
eks-5308cf7

# Our kubectl context should be automatically changed to new cluster
kubectl config view --minify
```

## kubectl commands

```
kubectl get pods -n kube-system
kubectl get pods -n kube-system -o wide
```

- Create the Deployment by running the following commands:
  - Reference : [run-stateless-application-deployment](#)

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml
wget https://k8s.io/examples/controllers/nginx-deployment.yaml

[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get events
LAST SEEN   TYPE      REASON          OBJECT
```

## MESSAGE

```
2m48s      Normal    Scheduled    pod/nginx-deployment-66b6c48dd5-qkst9
Successfully assigned default/nginx-deployment-66b6c48dd5-qkst9 to ip-192-168-21-
139.ap-south-1.compute.internal
2m47s      Normal    Pulling     pod/nginx-deployment-66b6c48dd5-qkst9
Pulling image "nginx:1.14.2"
2m38s      Normal    Pulled      pod/nginx-deployment-66b6c48dd5-qkst9
Successfully pulled image "nginx:1.14.2" in 9.12123707s
2m37s      Normal    Created     pod/nginx-deployment-66b6c48dd5-qkst9
Created container nginx
2m37s      Normal    Started     pod/nginx-deployment-66b6c48dd5-qkst9
Started container nginx
37s        Warning   FailedScheduling pod/nginx-deployment-66b6c48dd5-x92pj
0/2 nodes are available: 2 Too many pods.
2m48s      Normal    Scheduled    pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully assigned default/nginx-deployment-66b6c48dd5-zb9bb to ip-192-168-46-
56.ap-south-1.compute.internal
2m47s      Normal    Pulling     pod/nginx-deployment-66b6c48dd5-zb9bb
Pulling image "nginx:1.14.2"
2m37s      Normal    Pulled      pod/nginx-deployment-66b6c48dd5-zb9bb
Successfully pulled image "nginx:1.14.2" in 9.470360458s
2m37s      Normal    Created     pod/nginx-deployment-66b6c48dd5-zb9bb
Created container nginx
2m37s      Normal    Started     pod/nginx-deployment-66b6c48dd5-zb9bb
Started container nginx
2m48s      Normal    SuccessfulCreate replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-zb9bb
2m48s      Normal    SuccessfulCreate replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-qkst9
2m48s      Normal    SuccessfulCreate replicaset/nginx-deployment-66b6c48dd5
Created pod: nginx-deployment-66b6c48dd5-x92pj
2m48s      Normal    ScalingReplicaSet deployment/nginx-deployment
Scaled up replica set nginx-deployment-66b6c48dd5 to 3
```

```
kubectl describe pods -n default > pods.txt
cat pods.txt
```

```
Name:      nginx-deployment-66b6c48dd5-qkst9
Namespace: default
Priority:   0
Node:      ip-192-168-21-139.ap-south-1.compute.internal/192.168.21.139
Start Time: Wed, 02 Feb 2022 19:29:23 +0000
Labels:    app=nginx
           pod-template-hash=66b6c48dd5
Annotations: kubernetes.io/psp: eks.privileged
Status:    Running
IP:        192.168.15.247
IPs:
  IP:      192.168.15.247
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
  nginx:
    Container ID:
      docker://cb8aaf025035aa3a81d05f1598efbfc945bbe4ae01558a0ee0dee02942f32175
```

```

    Image:          nginx:1.14.2
    Image ID:       docker-
pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160
306b8d
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
        Started:   Wed, 02 Feb 2022 19:29:34 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
        /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-5qrpp
(ro)
Conditions:
  Type            Status
  Initialized      True
  Ready           True
  ContainersReady True
  PodScheduled    True
Volumes:
  kube-api-access-5qrpp:
    Type:          Projected (a volume that contains injected data from
multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute op=Exists for
300s
                  node.kubernetes.io/unreachable:NoExecute op=Exists
for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   9m3s  default-scheduler  Successfully assigned
default/nginx-deployment-66b6c48dd5-qkst9 to ip-192-168-21-139.ap-south-
1.compute.internal
  Normal  Pulling     9m2s  kubelet        Pulling image "nginx:1.14.2"
  Normal  Pulled      8m53s  kubelet        Successfully pulled image
"nginx:1.14.2" in 9.12123707s
  Normal  Created     8m52s  kubelet        Created container nginx
  Normal  Started     8m52s  kubelet        Started container nginx

Name:          nginx-deployment-66b6c48dd5-x92pj
Namespace:     default
Priority:       0
Node:          <none>
Labels:        app=nginx
                pod-template-hash=66b6c48dd5
Annotations:   kubernetes.io/psp: eks.privileged
Status:        Pending

```

```

IP:
IPs:          <none>
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
  nginx:
    Image:      nginx:1.14.2
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-x7mxn
(ro)
Conditions:
  Type          Status
  PodScheduled  False
Volumes:
  kube-api-access-x7mxn:
    Type:          Projected (a volume that contains injected data from
multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute op=Exists for
300s
                 node.kubernetes.io/unreachable:NoExecute op=Exists
for 300s
Events:
  Type          Reason          Age          From          Message
  ----          -
  Warning       FailedScheduling  52s (x9 over 9m3s)  default-scheduler  0/2 nodes are
available: 2 Too many pods.

Name:          nginx-deployment-66b6c48dd5-zb9bb
Namespace:     default
Priority:       0
Node:          ip-192-168-46-56.ap-south-1.compute.internal/192.168.46.56
Start Time:    Wed, 02 Feb 2022 19:29:23 +0000
Labels:        app=nginx
               pod-template-hash=66b6c48dd5
Annotations:   kubernetes.io/psp: eks.privileged
Status:        Running
IP:            192.168.51.252
IPs:
  IP:          192.168.51.252
Controlled By: ReplicaSet/nginx-deployment-66b6c48dd5
Containers:
  nginx:
    Container ID:
docker://9558a069050aba846aa62160509592ebb85dd82d1d90504bc5b03665a0259d47
    Image:      nginx:1.14.2
    Image ID:    docker-

```

```
pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160306b8d
```

```
Port:      80/TCP
Host Port: 0/TCP
State:     Running
  Started: Wed, 02 Feb 2022 19:29:34 +0000
Ready:     True
Restart Count: 0
Environment: <none>
Mounts:
```

```
/var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-tnsxj
```

(ro)

Conditions:

Type	Status
Initialized	True
Ready	True
ContainersReady	True
PodScheduled	True

Volumes:

kube-api-access-tnsxj:

Type: Projected (a volume that contains injected data from multiple sources)

TokenExpirationSeconds: 3607  
ConfigMapName: kube-root-ca.crt  
ConfigMapOptional: <nil>  
DownwardAPI: true

QoS Class: BestEffort

Node-Selectors: <none>

Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s

node.kubernetes.io/unreachable:NoExecute op=Exists

for 300s

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	9m3s	default-scheduler	Successfully assigned default/nginx-deployment-66b6c48dd5-zb9bb to ip-192-168-46-56.ap-south-1.compute.internal
Normal	Pulling	9m2s	kubelet	Pulling image "nginx:1.14.2"
Normal	Pulled	8m52s	kubelet	Successfully pulled image "nginx:1.14.2" in 9.470360458s
Normal	Created	8m52s	kubelet	Created container nginx
Normal	Started	8m52s	kubelet	Started container nginx

- The above command shows pods details where the pod gets an IP from ENI in the Subnet.

Network interfaces (6) <small>Info</small>									
<input type="text" value="Filter network interfaces"/>									
<input type="text" value="search:eks"/> <input type="button" value="Clear filters"/>									
Subnet ID	VPC ID	Availability Zone	Interface Type	Description	Instance ID	Status	Public IPv4...	Primary private IPv4...	Secondary private IPv4 add...
subnet-067941fe3a2fce5ab	vpc-02b76a6ce27d533a6	ap-south-1c	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.165.131	-
subnet-00f68a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	aws-K8S-I-08266ca7c...	i-08266ca7c22b644a5	In-use	-	192.168.30.4	192.168.15.247
subnet-00f68a720160ae9c2	vpc-02b76a6ce27d533a6	ap-south-1a	Elastic network interface	-	i-08266ca7c22b644a5	In-use	65.0.129.57	192.168.21.139	192.168.7.193
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	aws-K8S-I-0c544c004...	i-0c544c004ed1ea632	In-use	-	192.168.49.221	192.168.51.252
subnet-004a60aba384c11a	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	Amazon EKS test-eks-...	-	In-use	-	192.168.150.81	-
subnet-09aae16454bffc88b	vpc-02b76a6ce27d533a6	ap-south-1b	Elastic network interface	-	i-0c544c004ed1ea632	In-use	3.111.34.130	192.168.46.56	192.168.46.207

- AWS EKS supports native VPC networking with the Amazon VPC Container Network Interface (CNI) plugin for Kubernetes.
- Using this plugin allows Kubernetes Pods to have the same IP address inside the pod as they do on the VPC network.
- For more information, see [amazon-vpc-cni-k8s](#) and [Proposal: CNI plugin for Kubernetes networking over AWS VPC](#) on GitHub.
- The Amazon VPC CNI plugin is fully supported for use on Amazon EKS and self-managed Kubernetes clusters on AWS.
- Refer the ENI Limit as per Instance Type :  
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html#AvailableIpPerENI>
  - The formula for defining the maximum number of Pods per EC2 Node instance is as follows:
    - $N * (M-1) + 2$ 
      - $N$  is the number of Elastic Network Interfaces (ENI) of the instance type
      - $M$  is the number of IP addresses per ENI.
    - For e.g. **t2.micro** instance, this calculation is  $2 * (2-1) + 2 = 4$  Pods
- Use below command to get details on same:

```
aws ec2 describe-instance-types --filters "Name=instance-type,Values=t4g.*" --
query "InstanceTypes[.]{Type: InstanceType, MaxENI:
NetworkInfo.MaximumNetworkInterfaces, IPv4addr:
NetworkInfo.Ipv4AddressesPerInterface}" --output table
```

DescribeInstanceTypes			
IPv4addr	MaxENI	Type	
2	2	t4g.nano	
6	3	t4g.medium	
12	3	t4g.large	
15	4	t4g.2xlarge	
4	3	t4g.small	
2	2	t4g.micro	
15	4	t4g.xlarge	

```
kubectl get pods --all-namespaces -o wide
kubectl get pods --all-namespaces -o wide | grep -i running
[cloudshell-user@ip-10-0-50-112 ~]$ kubectl get pods --all-namespaces -o wide |
grep -i running
default          nginx-deployment-66b6c48dd5-qkst9  1/1      Running   0   38m
192.168.15.247 ip-192-168-21-139.ap-south-1.compute.internal
default          nginx-deployment-66b6c48dd5-zb9bb  1/1      Running   0   38m
192.168.51.252 ip-192-168-46-56.ap-south-1.compute.internal
kube-system      aws-node-jls58                      1/1      Running   0   125m
192.168.46.56 ip-192-168-46-56.ap-south-1.compute.internal
kube-system      aws-node-tmw12                      1/1      Running   0   125m
```



```

192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system   coredns-7f95bc96cc-dl6dn          1/1    Running   0    135m
192.168.46.207 ip-192-168-46-56.ap-south-1.compute.internal
kube-system   coredns-7f95bc96cc-fj56b          1/1    Running   0    135m
192.168.7.193  ip-192-168-21-139.ap-south-1.compute.internal
kube-system   kube-proxy-dqlph                   1/1    Running   0    125m
192.168.21.139 ip-192-168-21-139.ap-south-1.compute.internal
kube-system   kube-proxy-xlpfk                   1/1    Running   0    125m
192.168.46.56  ip-192-168-46-56.ap-south-1.compute.internal
[cloudshell-user@ip-10-0-50-112 ~]$

```

- Delete the cluster using

```
eksctl delete cluster --name test-eks-cluster
```

## Reference

- [EKS Cost Optimization](#)
- [Kubernetes instance calculator](#)
- [Resource Limits and Requests](#)

## EKS Cluster Pricing

- EKS is not free (Unlike other AWS Services), In short, no free-tier for EKS.

### EKS Cluster Pricing

- Below is the EKS Cluster Pricing:
  - \$0.10/hour
  - \$2.4/day
  - \$72/month

### EKS Worker Nodes Pricing-EC2

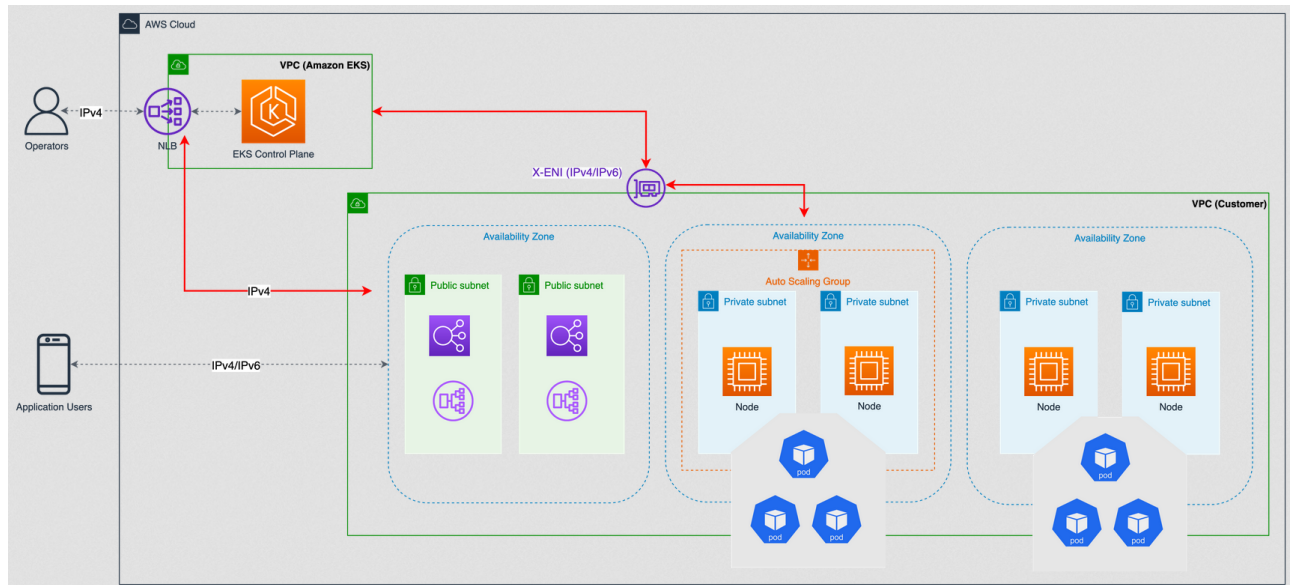
- You pay for AWS resources (e.g. EC2 instances or EBS volumes)
- T3 Medium Server in N.Virginia
  - \$0.0416 per Hour
  - Per Day: \$0.9984 - Approximately \$1
  - Per Month: \$30 per 1 t3.medium server
- Reference: [EC2 On-Demand Pricing](#)

## EKS Case Study:

- [meity-gov-india-case-study](#)
- [Architectural lessons from CoWIN platform](#)
- [cowin-public-v2](#)

## Reference:

- [AWS Graviton2 t4g.small free for all AWS Accounts until December 31, 2023](#)
  - Until December 31, 2023, all AWS customers will be enrolled automatically in the t4g.small Free Trial as detailed in the AWS Free Tier.
  - During the free-trial period, customers who run a t4g.small instance will automatically get 750 free hours per month deducted from their bill during each month.
- [aws-eks-best-practices](#)



- [eks-cluster-connection](#)
- [terraform-eks](#)