



# Kubernetes : Overview

- What is Kubernetes?
- Kubernetes Nodes
- Kubernetes Architecture:
- Kubernetes Components:
- Running and managing containers using Kubernetes
- Kubernetes Cluster
- Application Deployment Model
- Amazon EKS
- Companies adopting Amazon EKS
- Kubernetes Provides :
- K8s : Customer Case Studies

# What is Kubernetes?

- **Docker** is a container runtime while **Kubernetes** is a platform for running and managing containers.
- It is an open source **orchestration platform** for automating **deployment, scaling and the operations** of **application containers** across *clusters of hosts*.
- It is also defined as a platform for **creating, deploying and managing various distributed applications**.
- These applications may be of different sizes and shapes.
- **Kubernetes was originally developed by Google** to **deploy scalable, reliable systems in containers** via application-oriented APIs.

# Kubernetes Nodes

- Kubernetes follows **master-worker** architecture.
- Kubernetes architecture has a **master node** and **worker nodes**.

There are four components of a **master node**.

- Kube API Server
- controller
- scheduler
- etcd

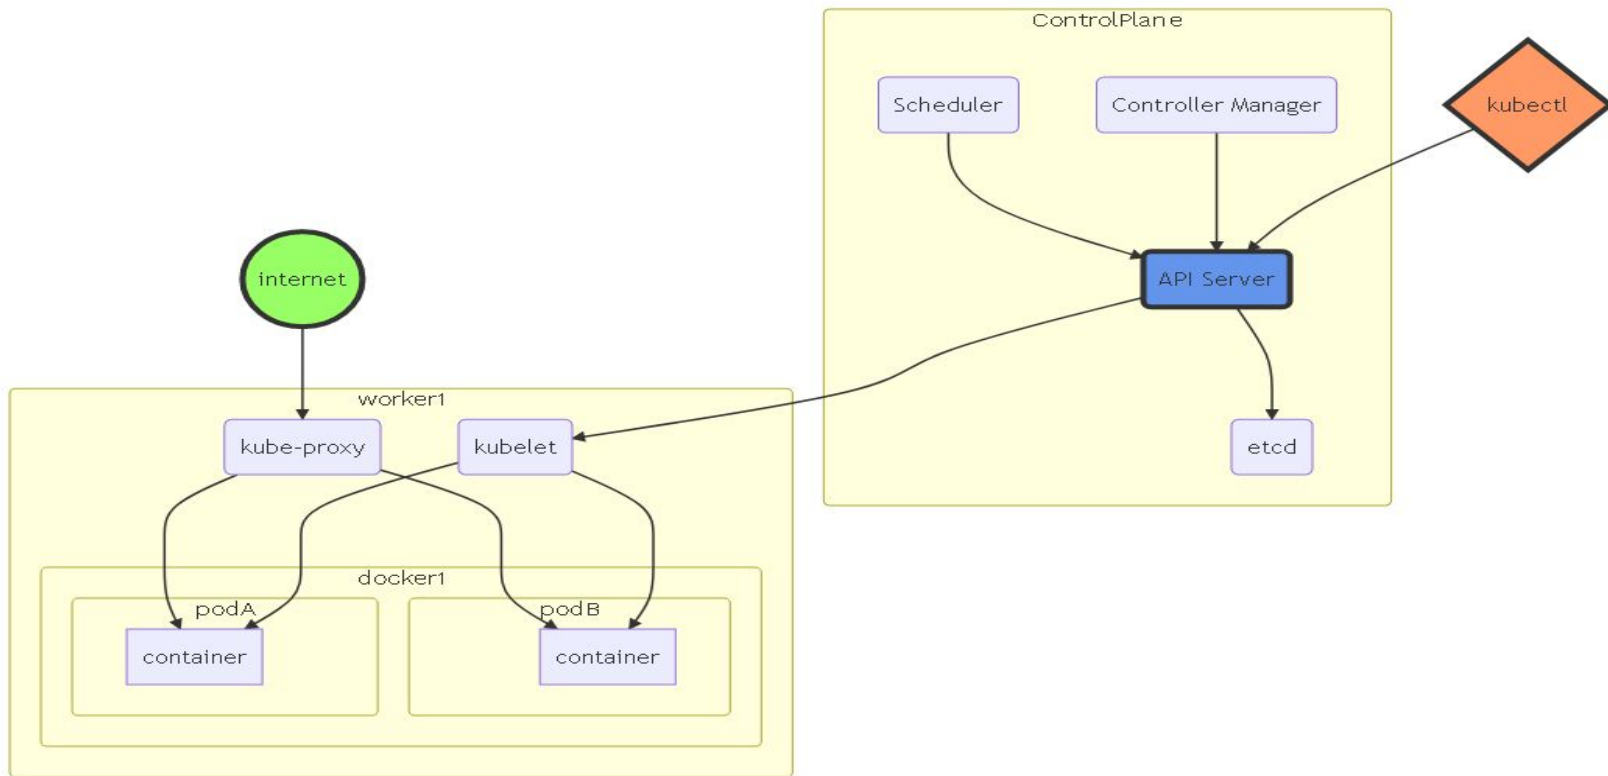
There are three components of a **worker node**.

- kubelet
- kube-proxy
- Container runtime


A worker node is a **virtual or physical server** that runs the applications and is controlled by the master node.

Each node has a **Kubelet**, which is an agent for managing the node and communicating with the Kubernetes control plane.

# Kubernetes Architecture:



# Master Node



- **API Server** performs all the administrative tasks on the master node. A **user sends the rest commands to the API server**, which then validates the requests, then processes and executes them.
- The **scheduler**, schedules the work to different worker nodes. It has the **resource usage information** for each worker node. The scheduler considers service requirements and **node resource parameters** and schedules the work in terms of **pods and services**.
- The **controller manager** makes sure that your **current state is the same as the desired state**. A **controller** is a loop that continually monitors your cluster and performs actions when certain events occur. The **controller manager** oversees all the controllers in your cluster. It starts their processes and ensures they're operational the whole time that your cluster's running.
- The **etcd** is a **key-value database store** that stores the state of the **cluster,node or pods**. It is also used to store the **configuration** details such as the **ConfigMaps** ( store environment variables ) and **Secrets**
  - **etcd** saves the resulting state of the cluster as a **key-value** store.

# Worker Node

## Worker Node

Kubelet


Kube-Proxy

Container Runtime  
(Docker)


- **Kubelet** is basically an **agent/daemon** that runs on **each worker node** and **communicates with the master node**. This agent is responsible for making sure that **containers are running in a Pod on a node**.
- The **container runtime** is basically used to run and manage the continuous life cycle of the container on the worker node.
- **Kube-proxy** runs on each worker node as the **network proxy**. It listens to the API server for each **service** point creation or deletion. For each service point, **kube-proxy** sets the routes so that it can reach to it.



# Worker Node



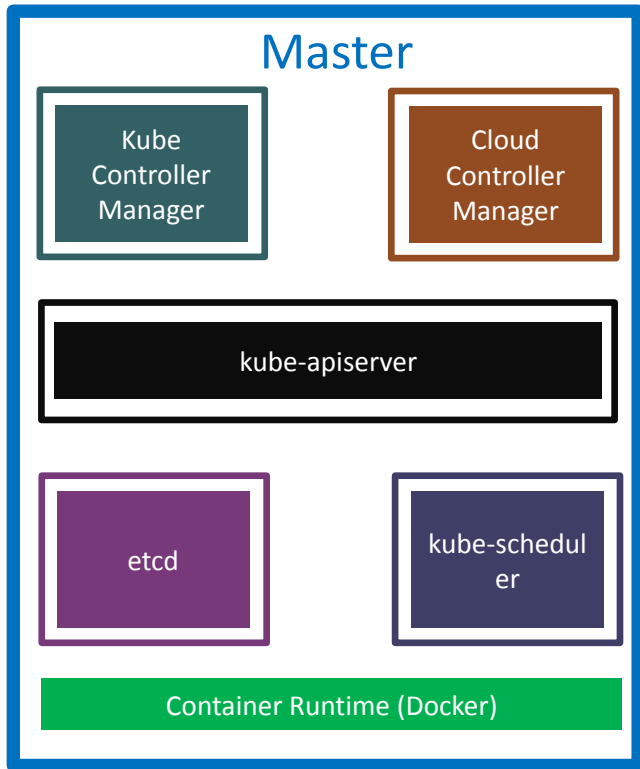


# Kubernetes Architecture:





# Kubernetes Architecture:



# Kubernetes Objects/Components:

**Pods**: Co-located group of containers that share an **IP, namespace, storage volume**.

A pod is a collection of **application containers** and volumes running in the same **execution environment**. Most of the **pod manifests** are written using **YAML** or JSON scripts

**Replica Sets**: maintains a stable set of replica Pods running at any given time.

**Manages the lifecycle of pods, and If any Pod ( Container ) goes down, RS launches another Pod.**

**Deployment**: A **Deployment runs multiple replicas of your application containers** and automatically replaces any **containers** that fail or become unresponsive. **Rollout & rollback images changes to applications**. Deployments are well-suited for **stateless** applications.

**Label**: **Key/Value pairs used for association and filtering**. Also select group of objects ( **Tags** ), **Selectors** is used to select specific K8s Object using the Label key:values.

**Service**: Maps a fixed IP address to a logical group of pods.

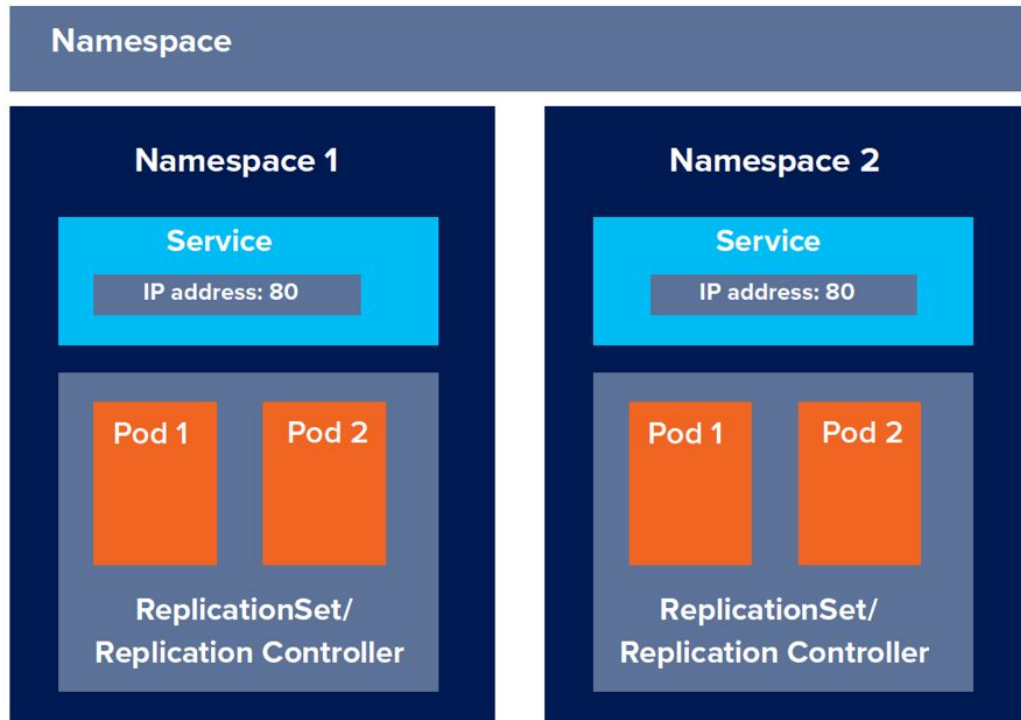
**Ingress**: These are objects that provide an easy-to-use front-end (externalised API surface area).

# Kubernetes Components:

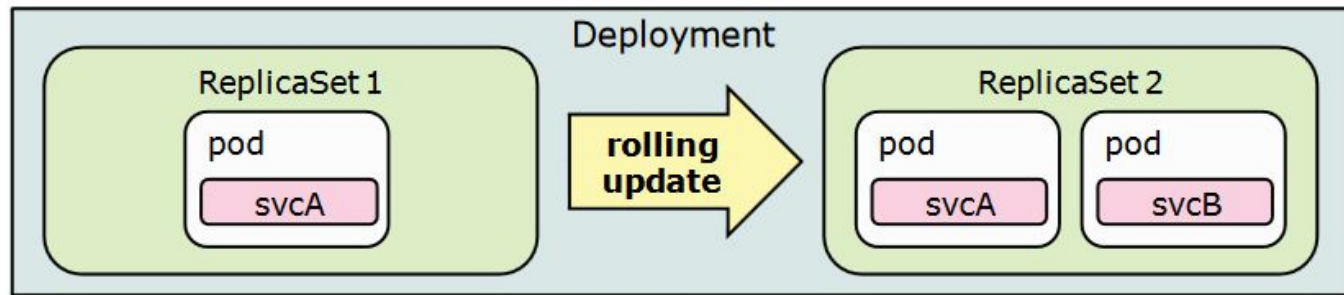
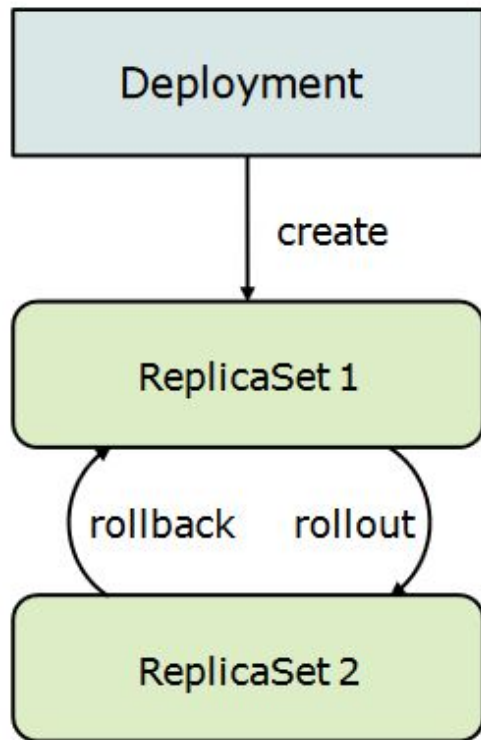
**NameSpaces**: It is used to remove name collision within a cluster. It supports **multiple virtual clusters** on the same **physical cluster**. This provides isolation and complete access to control the degree to which other services interact with it. ( **default** namespace ). Namespaces provides a mechanism for isolating **groups of resources** within a single cluster.

- In a new cluster, Kubernetes automatically creates the following namespaces: **`default`** (for user workloads) and for the Kubernetes **control plane: `kube-system`**.
- **default**
  - All pods will be launched in this default namespace.
  - **kubectl get pods -n default**
- **kube-system ( do not make any changes to this namespace )**
  - All master components daemons will be running here
  - **kubectl get pods -n kube-system**

# Kubernetes Components:



# Kubernetes Components: Rolling update

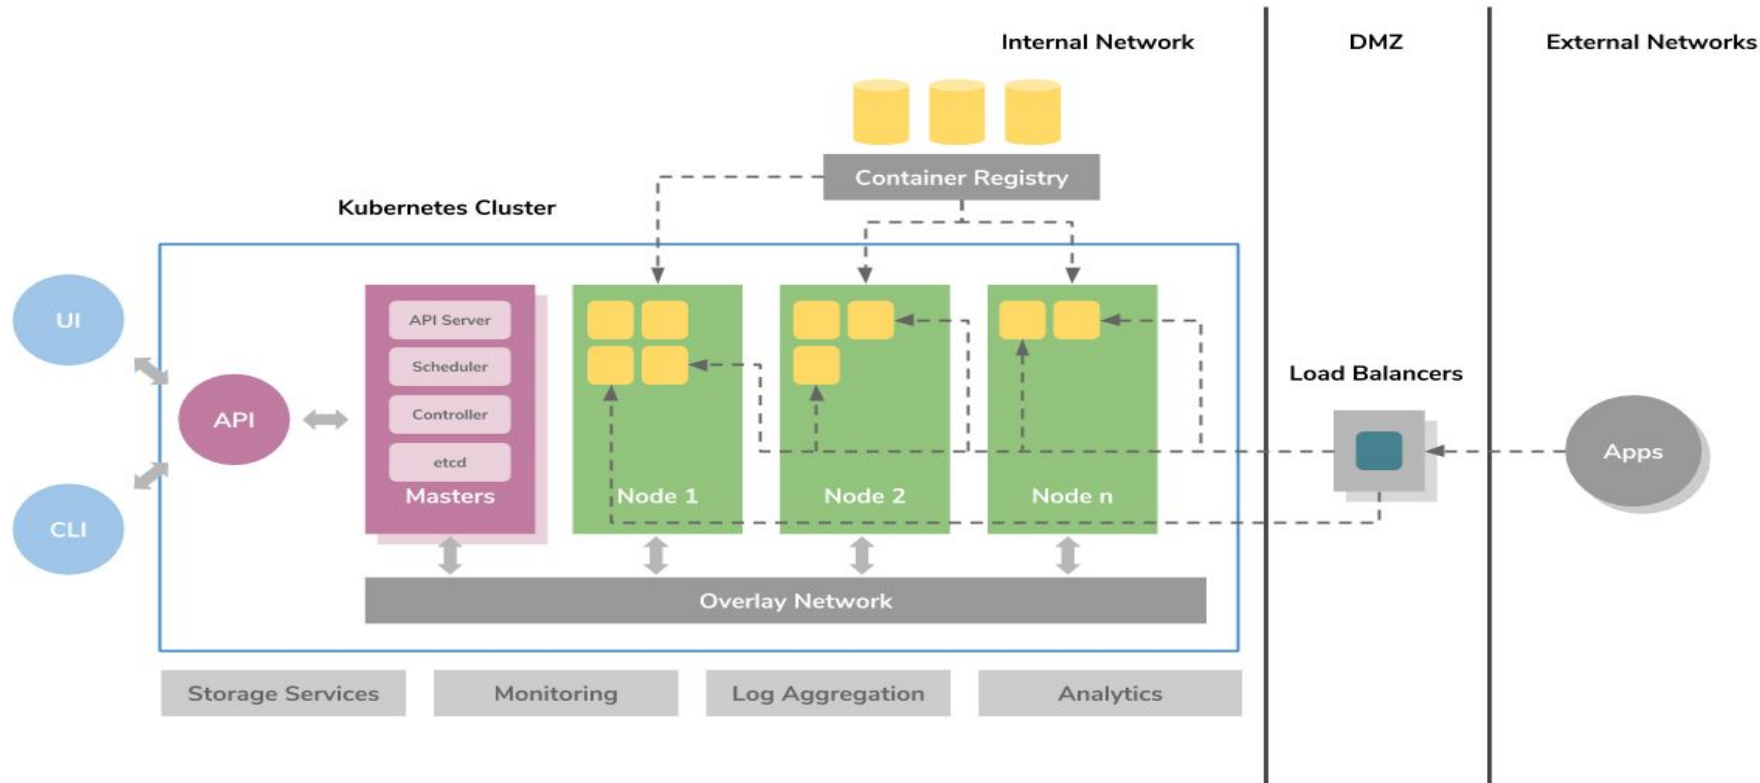




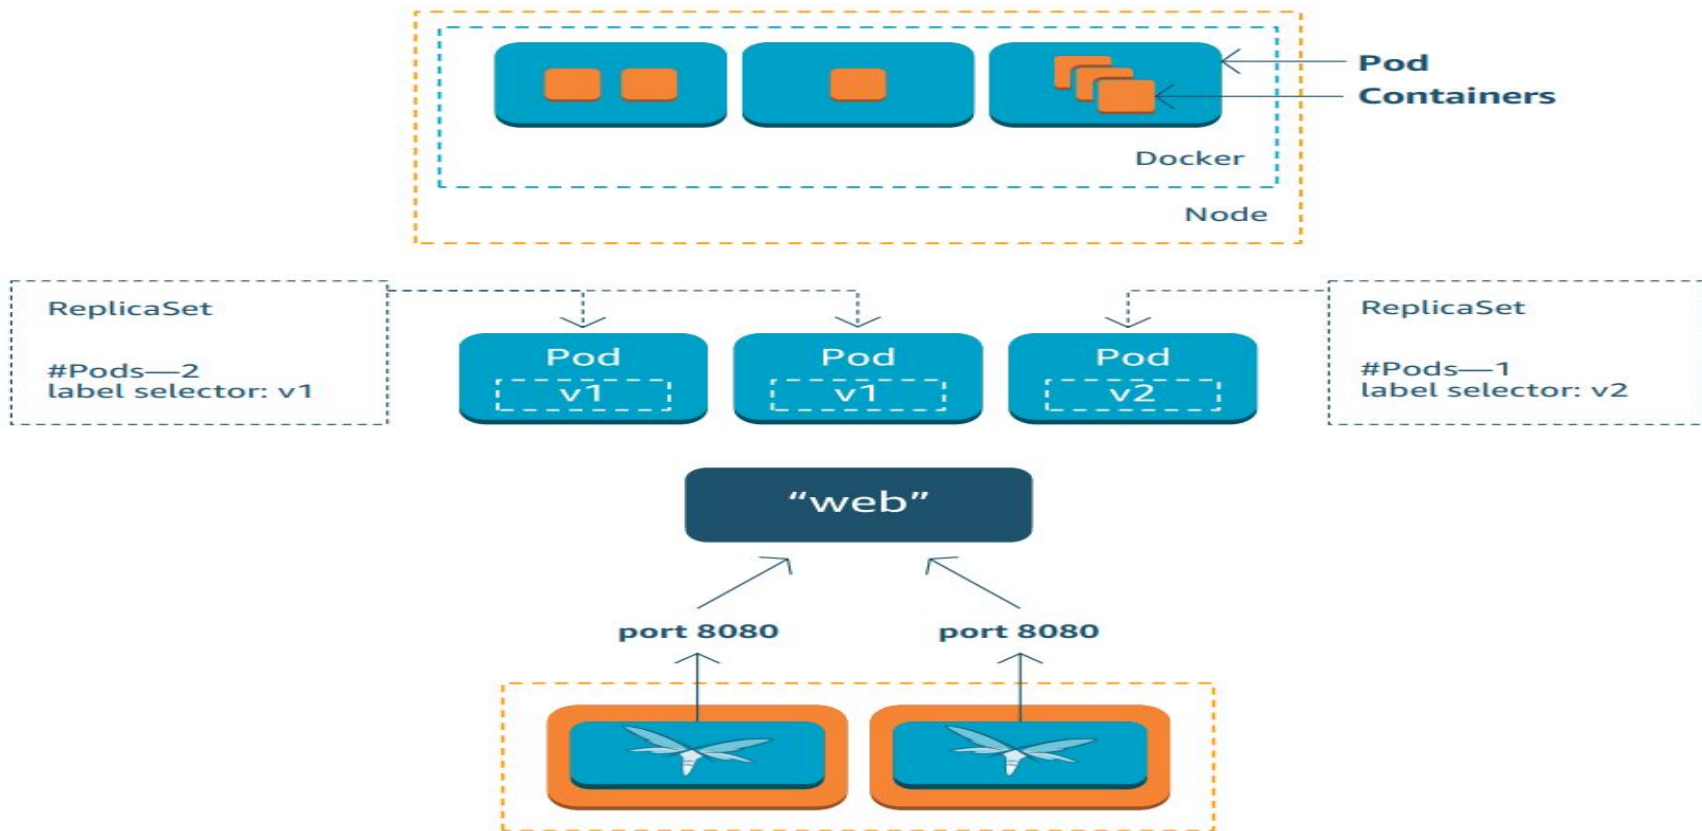
# Kubernetes Cluster

- **Kubernetes coordinates a highly available cluster of computers that are connected to work as a single unit.**
- A Kubernetes cluster consists of two types of resources:
  - The **Control Plane** coordinates the cluster
  - **Nodes** are the workers that run applications

# Kubernetes Architecture:

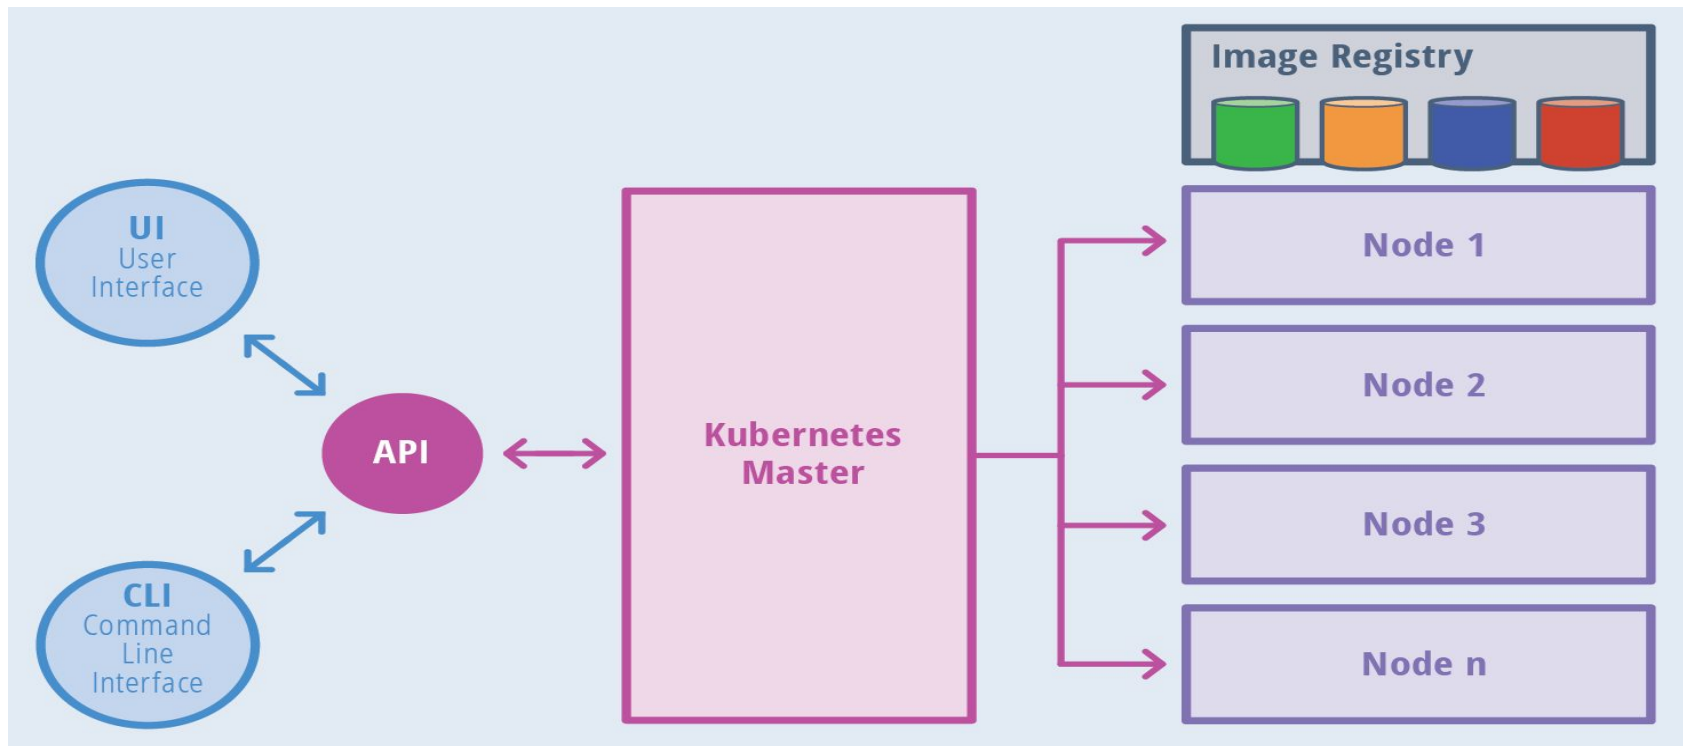


# Kubernetes Components:





# Kubernetes Architecture:



# Managing containers using Kubernetes

Docker provides many features by exposing the underlying **'cgroups'** technology provided by the Linux kernel. With this, the following resource usage can be managed and monitored:

- Kubernetes also can be installed using **Minikube**, locally.
- **Minikube** is a simulation of the Kubernetes cluster, but the main function of this is for **experimentation, local development or for learning purposes**.
- **Minikube is a single node kubernetes cluster**
- <https://kubernetes.io/docs/tutorials/hello-minikube/>
- <https://kubernetes.io/docs/tutorials/kubernetes-basics/create-cluster/cluster-intro/>
- Example Commands here:
- <https://kubernetes.io/docs/tasks/run-application/run-stateless-application-deployment/>
- <https://kubernetes.io/docs/concepts/workloads/controllers/deployment>

# Container Image

- A container image is a binary package that encapsulates all of the files necessary to run an application inside an OS container.
- The Open Image (OCI) is the standard image format that's most widely used. Container

## Two Types of container categories are:

- **System containers**, which try to imitate virtual machines and may run the full boot processes ( **kube-system** )
  - All master components are pods under this namespace.
- **Application containers**, which run single applications. ( **default** )
- System containers, which try to imitate virtual machines and may run the full boot processes.
- Application containers, which run single applications.
- The default container runtime used by Kubernetes is **Docker**.



# Kubernetes Service

- You can expose an application running on a set of PODs using different types of Services available in k8s.
- **ClusterIP**
  - Used for communication between applications inside k8s cluster
- **NodePort**
  - To access our application **outside of k8s cluster**, you can use **NodePort** service.
  - Exposes the Service on each **Worker Node's IP** at a static port (nothing but **NodePort**).
  - Port Range **30000-32767**
- **LoadBalancer**
  - Primarily for Cloud Providers to integrate with their Load Balancer services (Example: AWS Elastic Load Balancer)

# Kubernetes Architecture:


- Ability to deploy existing applications that run on VMs without any changes to the application code.
- On the high level, any application that runs on VMs can be deployed on Kubernetes by simply containerizing its components.
- Container grouping, container orchestration, overlay networking, container-to-container routing with layer 4 virtual IP based routing system, service discovery, support for running daemons, deploying stateful application components, and most importantly the ability to extend the container orchestrator for supporting complex orchestration requirements.

# Kubernetes Architecture:

- On very high level Kubernetes provides a set of dynamically scalable hosts for running workloads using containers and uses a set of management hosts called **masters** for providing an API for managing the entire container infrastructure.
- The workloads could include long-running services, batch jobs and container host specific daemons.
- All the container hosts are connected together using an overlay network for providing container-to-container routing.
- Applications deployed on Kubernetes are dynamically discoverable within the cluster network and can be exposed to the external networks using traditional load balancers.
- The state of the cluster manager is stored on a highly distributed key/value store etcd which runs within the master instances.



# Application Deployment Model



# Rolling Updates in K8s



Before

Live



V1.1

After

1 Node



OK?

2 Nodes



OK?

3 Nodes



4 Nodes



OK?

5 Nodes



OK?

6 Nodes



V1.2

Window Size = 1 node

A rolling deployment is a deployment strategy that updates running instances of an application with the new release. All nodes in a target environment are incrementally updated with the service or artifact version in integer N batches.





# Amazon EKS

- Amazon **Elastic Kubernetes Service** (Amazon EKS) makes it easy to deploy, manage, and scale containerized applications using [Kubernetes on AWS](https://aws.amazon.com/eks/).
- <https://aws.amazon.com/eks/>
  - <https://aws.amazon.com/eks/pricing/>
    - Pay \$0.10 for your master node
    - Pay underlying EC2 instance cost on the basis of type



# Companies adopting Amazon EKS



# Companies adopting Amazon EKS

- <https://www.infoworld.com/article/3664052/why-mercedes-benz-runs-on-900-kubernetes-clusters.html>
- <https://aws.amazon.com/solutions/case-studies/>
  - Filter by Domain to view all AWS Case Studies
- <https://aws.amazon.com/solutions/case-studies/itv-case-study/>
- <https://aws.amazon.com/solutions/case-studies/netflix/>



# DockerSwarm and Kubernetes

- Both Kubernetes and Docker Swarm are popular and used as **container orchestration platforms**.
- Docker Swarm is the native clustering for Docker.
- Originally, it did not provide much by way of container automation, but with the latest update to Docker Engine 1.12, container orchestration is now built into its core with first-party support.
- It takes some effort to get Kubernetes installed and running, as compared to the faster and easier Docker Swarm installation. Both have good scalability and high availability features built into them.
- Hence, one has to choose the right one based on the need of the hour.
- Do refer to <https://www.upcloud.com/blog/docker-swarm-vs-kubernetes/>



# Kubernetes Cluster

- **Kubernetes coordinates a highly available cluster of computers that are connected to work as a single unit.**
- A Kubernetes cluster consists of two types of resources:
  - The **Control Plane** coordinates the cluster
  - **Nodes** are the workers that run applications
- **Kubernetes Proxy** - for routing network traffic for load balancing services  
(<https://kubernetes.io/docs/getting-started-guides/scratch/>)
- **Kubernetes DNS** - a DNS server for naming and discovery of the services that are defined in DNS
- **Kubernetes UI** - this is the GUI to manage the clusters

# Managing containers using Kubernetes

- **Memory resources management and limitation**
- **CPU resources management and limitation**
- Running a Stateless application
  - <https://kubernetes.io/docs/tasks/run-application/run-stateless-application-deployment/>
- Running a Stateful application
  - <https://kubernetes.io/docs/tasks/run-application/run-single-instance-stateful-application/>